



University of
Pittsburgh

Applied Cryptography and Network Security

CS 1653



Summer 2023
Sherif Khattab
ksm73@pitt.edu

(Slides are adapted from Prof. Adam Lee's CS1653 slides.)

Announcements

- We are losing **five classes** this term
 - Two holidays are on a Monday
 - Midterm and final exams in class
 - I have to miss class on Monday June 12
- Two polls:
 - add 15 minutes to each lecture starting next week?
 - updated lecture time: 1:30 pm – 3:30 pm
 - makeup for ~ 2 lectures
 - agree on a time for two extra lectures outside regular class time

Please mark all your available times



Computer security is typically defined with respect to three types of properties

How do I ensure that my secrets remain secret?




Confidentiality

Can I trust the services that I use?



Integrity



Availability

Am I able to do what I need to do?

How do we characterize insecurity?

A **threat** is a **potential** violation of security. That is, a threat is something that you want to prevent from happening.

Note: A violation **need not occur** in order for a threat to exist!

Threats are sometimes broadly classified into four categories:

Disclosure: Information leakage

Deception: Acceptance of false information

Disruption: Interruption or prevention of correct operation

Usurpation: Unauthorized control of some part of a system

These four classes encompass many, many threats...

A few examples...

Snooping: Unauthorized interception of information

Category: Disclosure

Defenses: Confidentiality mechanisms

Modification or Alteration: Unauthorized change of information

Category: Deception, Disruption, and/or Usurpation

Defenses: Integrity mechanisms

Spoofing or Masquerading: Impersonation of one entity by another

Category: Deception and/or Usurpation

Defenses: Integrity mechanisms (e.g., crypto or authentication services)

Threats simply define the types of insecurity that we are concerned with...

Vulnerabilities are situations or conditions that allow a threat to be realized

For example:

Vulnerabilities in Design

Incorrect assumptions about operating environment

Vulnerabilities in Implementation

Double free bug

Backdoors (intentional or accidental)

Corrupted, Faulty or Malicious hardware (e.g., keylogging keyboards)

Vulnerabilities in Configuration

Incompletely specified policies (e.g., in firewalls)

There are many online databases, such as CERT/CC, that list vulnerabilities in widely-used software packages


An **attack** is the **exploitation** of a vulnerability to realize a threat

An Example: Bob's Bank

Bob's Bank runs an **online banking site** that allows its users to check their account balances and pay bills online

Clients can only see the data in their own account

Due to financial woes, the banking portal is run using an out-of-date web server



Threat: Seeing another user's bank records

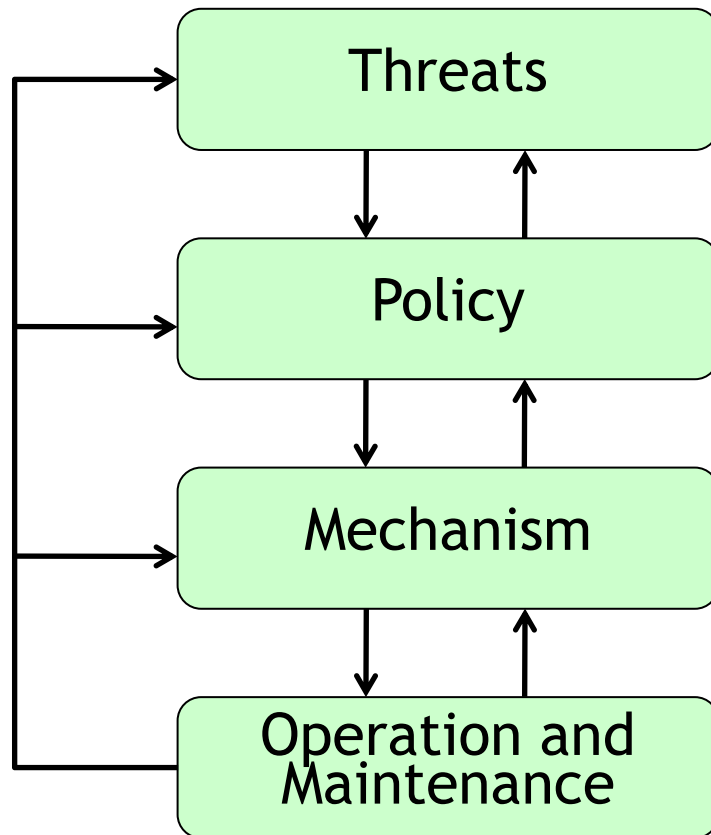
Vulnerability: Unsafe string operations in C code of server (e.g., strcpy)

Attack: A buffer overflow that gives the adversary root-level access to the system

So what?

**Security is not an absolute
property!**

Security is a process!



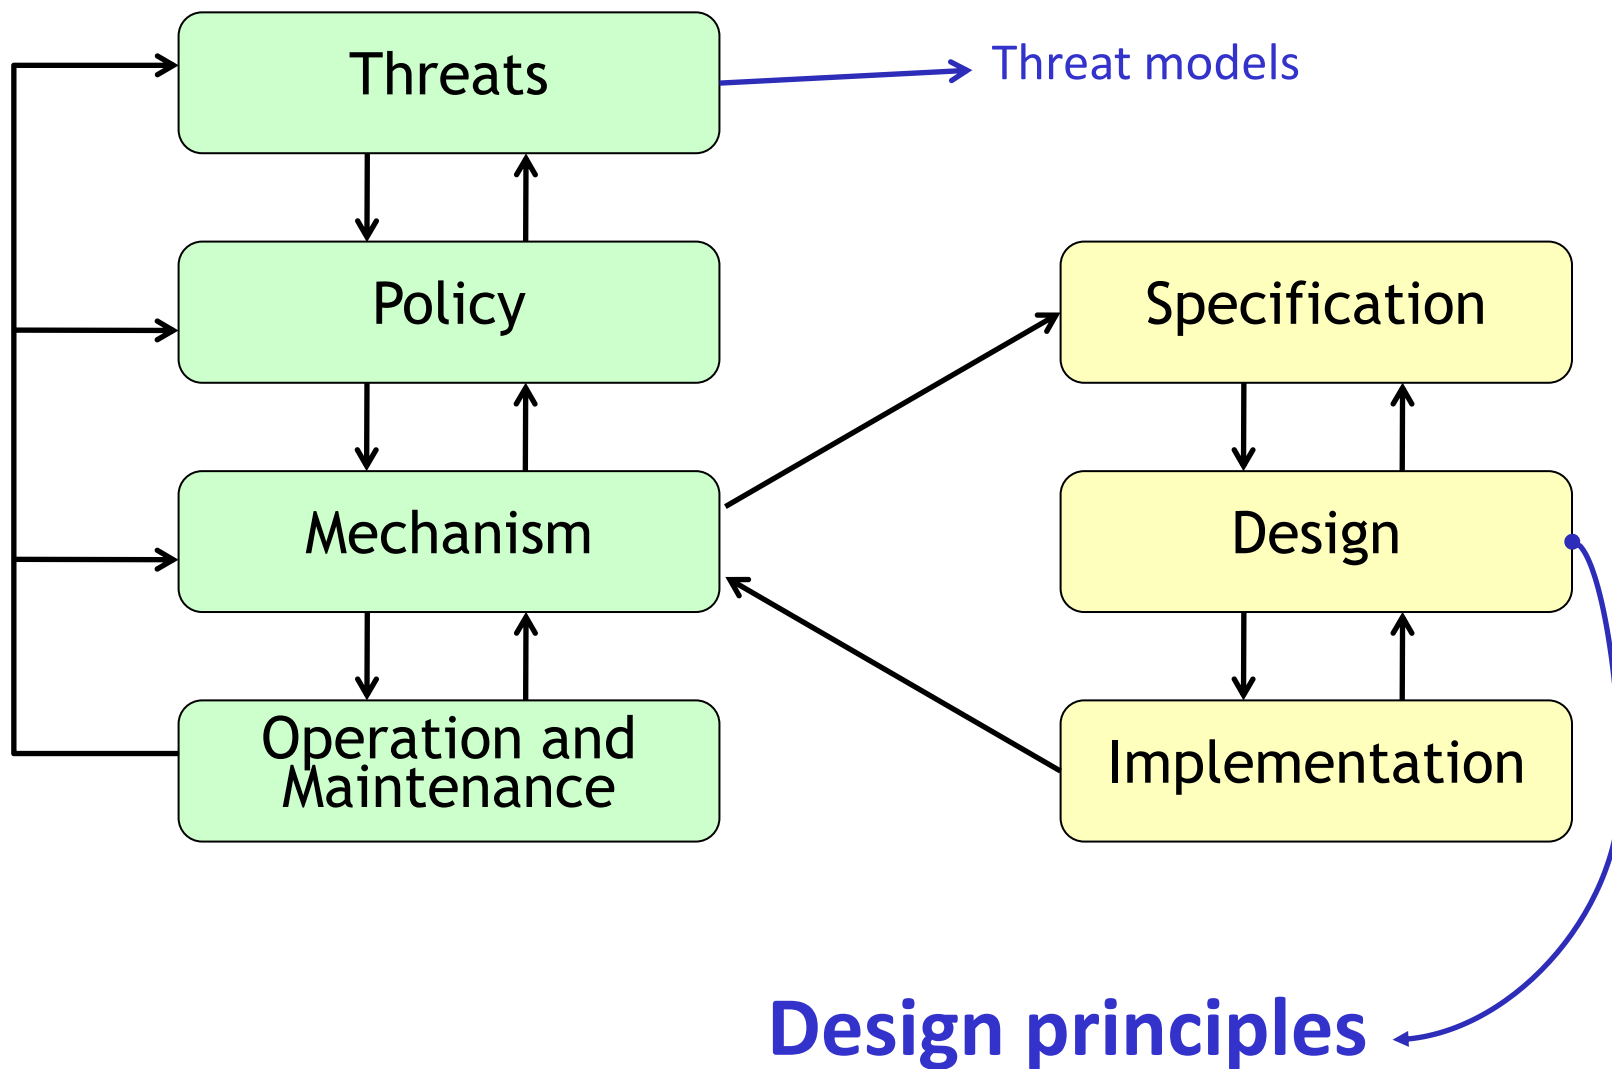
Steps:

1. Identify threats for the domain of interest
2. Define policies to protect against these threats
 - High-level organizational policies
 - Low-level logical policies
 - Everything in between!
3. Develop mechanisms to enforce these policies
4. Wash, rinse, repeat

Discussion

- What are your opinions regarding **online vulnerability databases** and/or bug bounty programs?
- Do they represent a **helpful tool** to the community, or simply provide attackers with an easy means of exploiting systems?
- What do you feel should be the “**protocol**” used regarding newly discovered vulnerabilities?

More details



Security is based on assumptions and trust

Some universities have **open exam** policies that allow students to work on their exams **at home (or anywhere else)** within a certain time window, as long as the student signs a statement indicating that the work in question is their own. Is this exam “system” secure?


This depends upon the assumptions made!

Trusted students → secure “system”

Malicious students → insecure “system”

Note: Violating this trust **assumption**
Invalidates the security of the system!


Is this OK?



Violating assumptions is a very easy way to discover vulnerabilities in computer systems


Many sources of assumptions

- Explicit threat models
- Software engineering documents
- Process descriptions or standards
- Insider information
- ...



Assumptions can also be implicit!

Example: Buffer overflows



Systematically define a good threat model before developing a system

Threat modeling can be a complicated software engineering process, a simple statement of assumptions, or anything in between

Software companies benefit from a process-based threat model

See “Threat Modeling” by Swiderski and Snyder (Microsoft Press, 2004)

In research papers, look for “Threat model”, “System model”, or “Environment” subsections that describe the assumptions made

Example: *We assume that a pub-sub system consists of a set of routing nodes as well as a trusted security manager node, while publishers and subscribers exist as applications outside of the pub-sub system. We assume that every publisher, subscriber, and routing node is managed by some principal in the set P of all principals. The security manager is a central authority that is trusted by publishers and subscribers to coordinate the system. Many existing pub-sub systems that protect publishers' private data require such a central authority for key management. Each principal $p_i \in P$ maintains a public key pair (K_i, K_i^{-1}) , and can obtain the public keys of other principals using a PKI or some other key distribution service. We assume that publishers publish data items from some value set V ...*

Fully specify ambiguous concepts!

In order for a threat model (or any specification) to be useful, it must be free of ambiguity

Example: Confusion over the term “risk”

One **entirely overloaded** word in computer security is “trust”

- Based on satisfying a concrete logical policy?

- Based upon subjective evaluation of past performance?

 - QoS based?

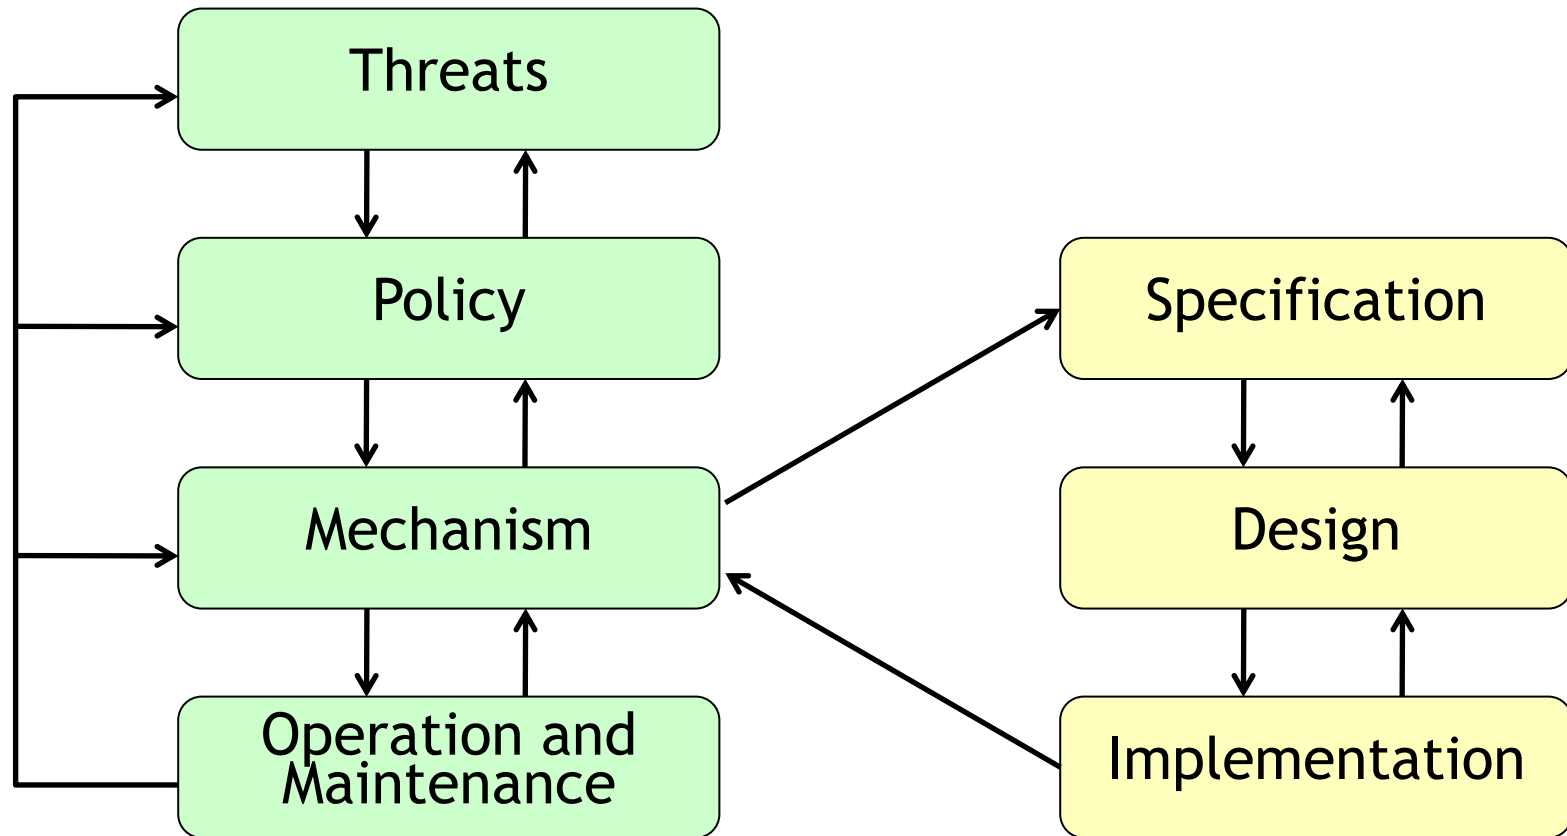
 - Attack based?

 - Selfishness/tragedy of the commons?

- Transitive trust? How does transitivity degrade trust?

- ...

Given an unambiguous threat model, policies can be defined and mechanisms designed



Security mechanisms are a means through which we develop assurance in a system

Assurance is an (often) approximate measure of how much a system can be trusted.

Assurance in a system can change over time

Increased belief in correctness over time (e.g., cryptography)

Decreased confidence after successful compromises (e.g., software)

Example: Drug Safety

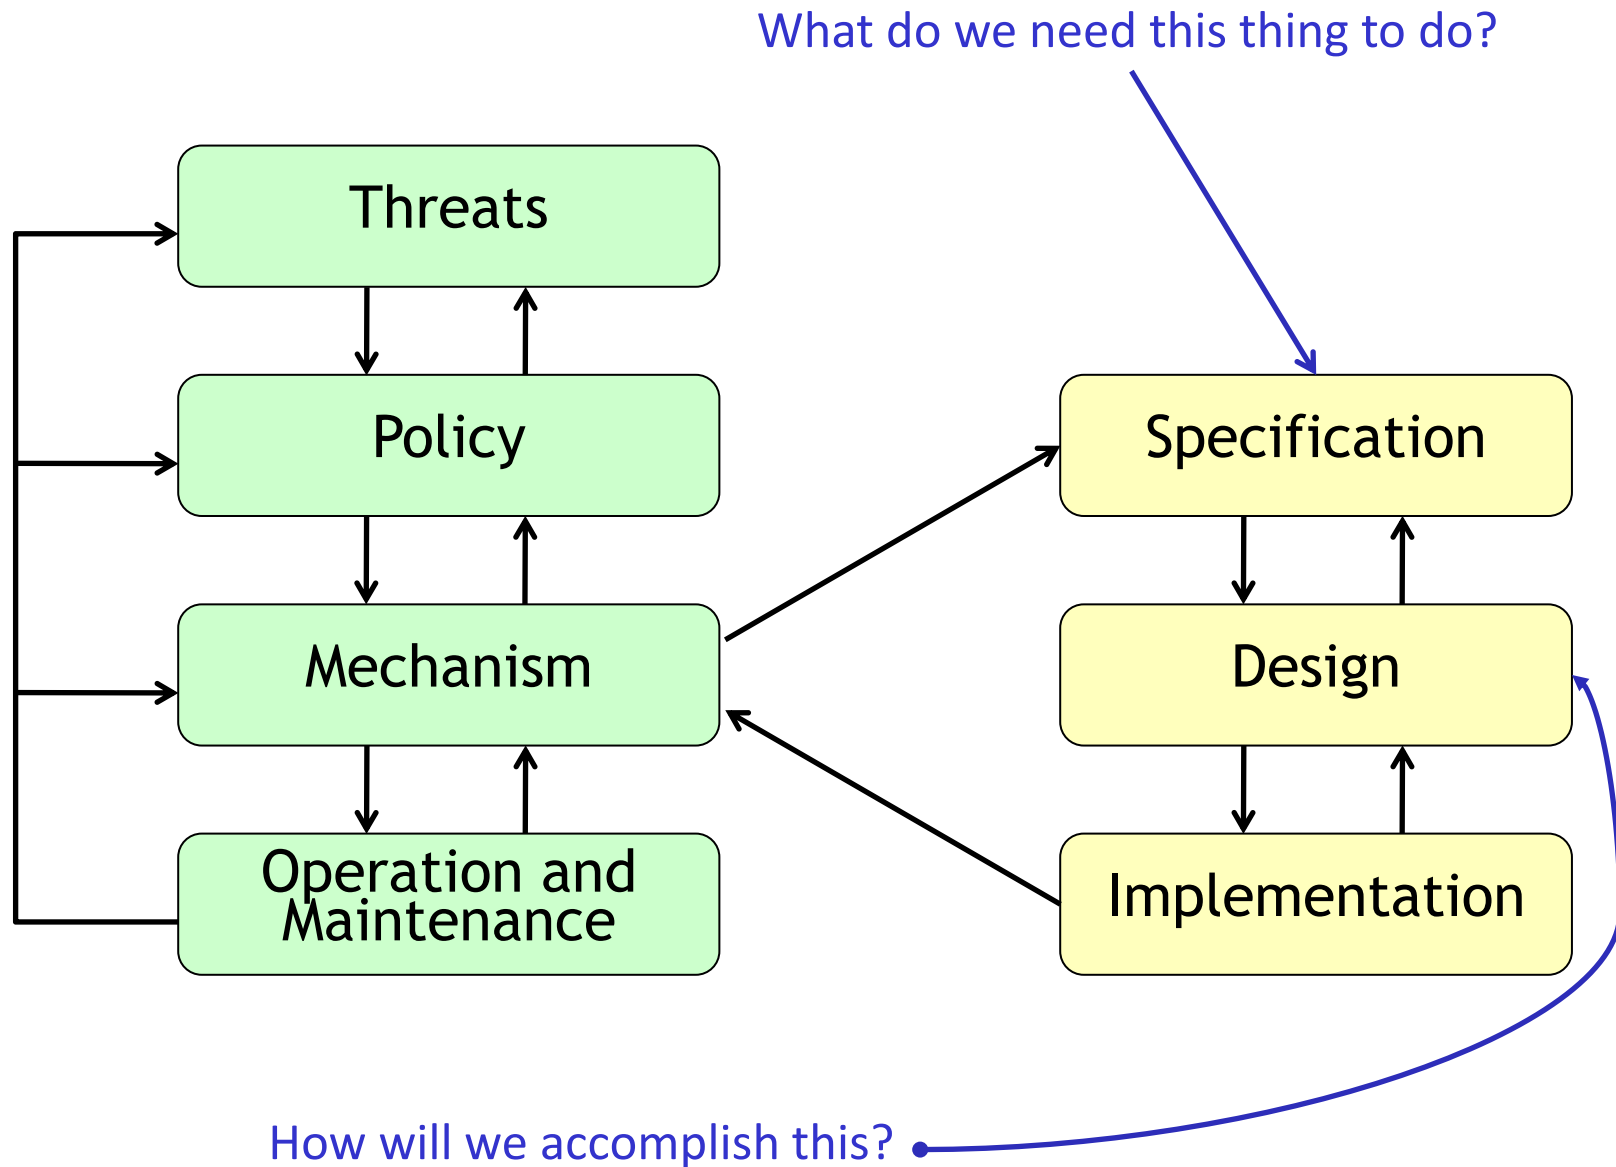
People typically believe that drugs manufactured in the US are safe because (i) the FDA enforces safety standards and (ii) the process control mechanisms used by companies ensure that drugs are not contaminated during manufacture.

What about contamination *after* manufacture?

Safety seals introduced after scares in the 80s.



Mechanism development should proceed in three steps



Specification

Definition: A **specification** is a collection of statements describing the desired functionality of a system. Specifications can be expressed in English, a formal logical language, or anything in between.

Specifications can be made at **any** level of detail.

Example: High Level

The computer should not be vulnerable to attack from the Internet.

How do you enforce this?!?

Example: Low Level

The computer should not accept any incoming network connections.

What about malicious content that is downloaded by the user?

Note: Specifications are used in many fields other than security.

Design

Definition: A system **design** translates a specification into components that will actually be implemented.

Example

Specification: The system shouldn't accept incoming network connections

Design: The system will contain a software firewall and will be located behind a network firewall. Both firewalls will prevent incoming connections.

Ultimately, it is important to show that a system design **satisfies** its specification. That is, the system should under no circumstance violate the conditions set forth in the specification.



Formal proof

How can we do this?

Informal
argumentation



Mechanism design should not be an ad-hoc process!

Invaluable reference:

Jerome H. Saltzer and Michael D. Schroeder, “The Protection of Information in Computer Systems,” Proceedings of the IEEE 63(9): 1278-1308, September 1975.

[Link on course web page](#)

See especially Section I, Part A(3)

Saltzer and Schroeder describe and justify the use of eight design principles for building secure and functional systems

Adopted with varying degrees of success over the years

Sometimes called the most-often cited and least-often read paper in computer security 😊

Principle (a): Economy of Mechanism

Definition: *Keep the design as simple and as small as possible.*

Why is this important for security?

Errors resulting in unwanted access probably won't be found during normal operation of a system

To check for these errors, line-by-line verification is important

This type of check is likely to fail with overly-complicated systems

This principle has many well-known incarnations

Philosophy: Occam's Razor

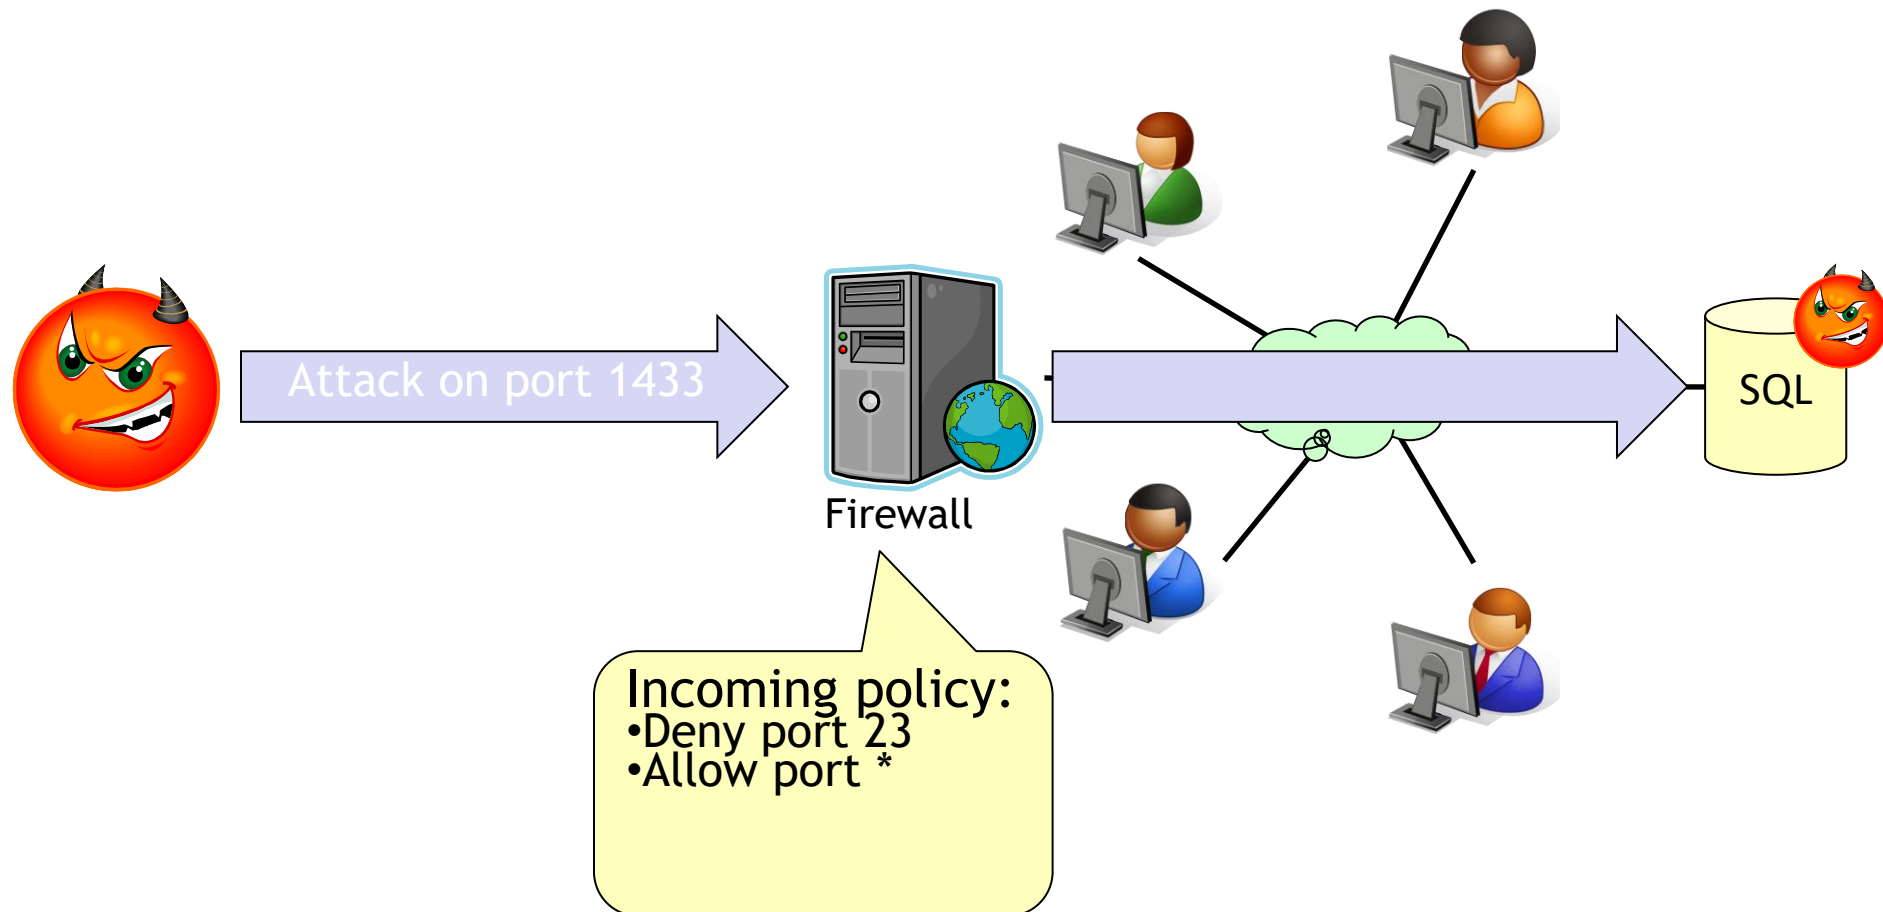
Everyday life: K.I.S.S.

Security practitioners are all too familiar with this principle...

Principle (b): Fail-Safe Defaults

Definition: Base decisions on *permissions* rather than *exclusions*.

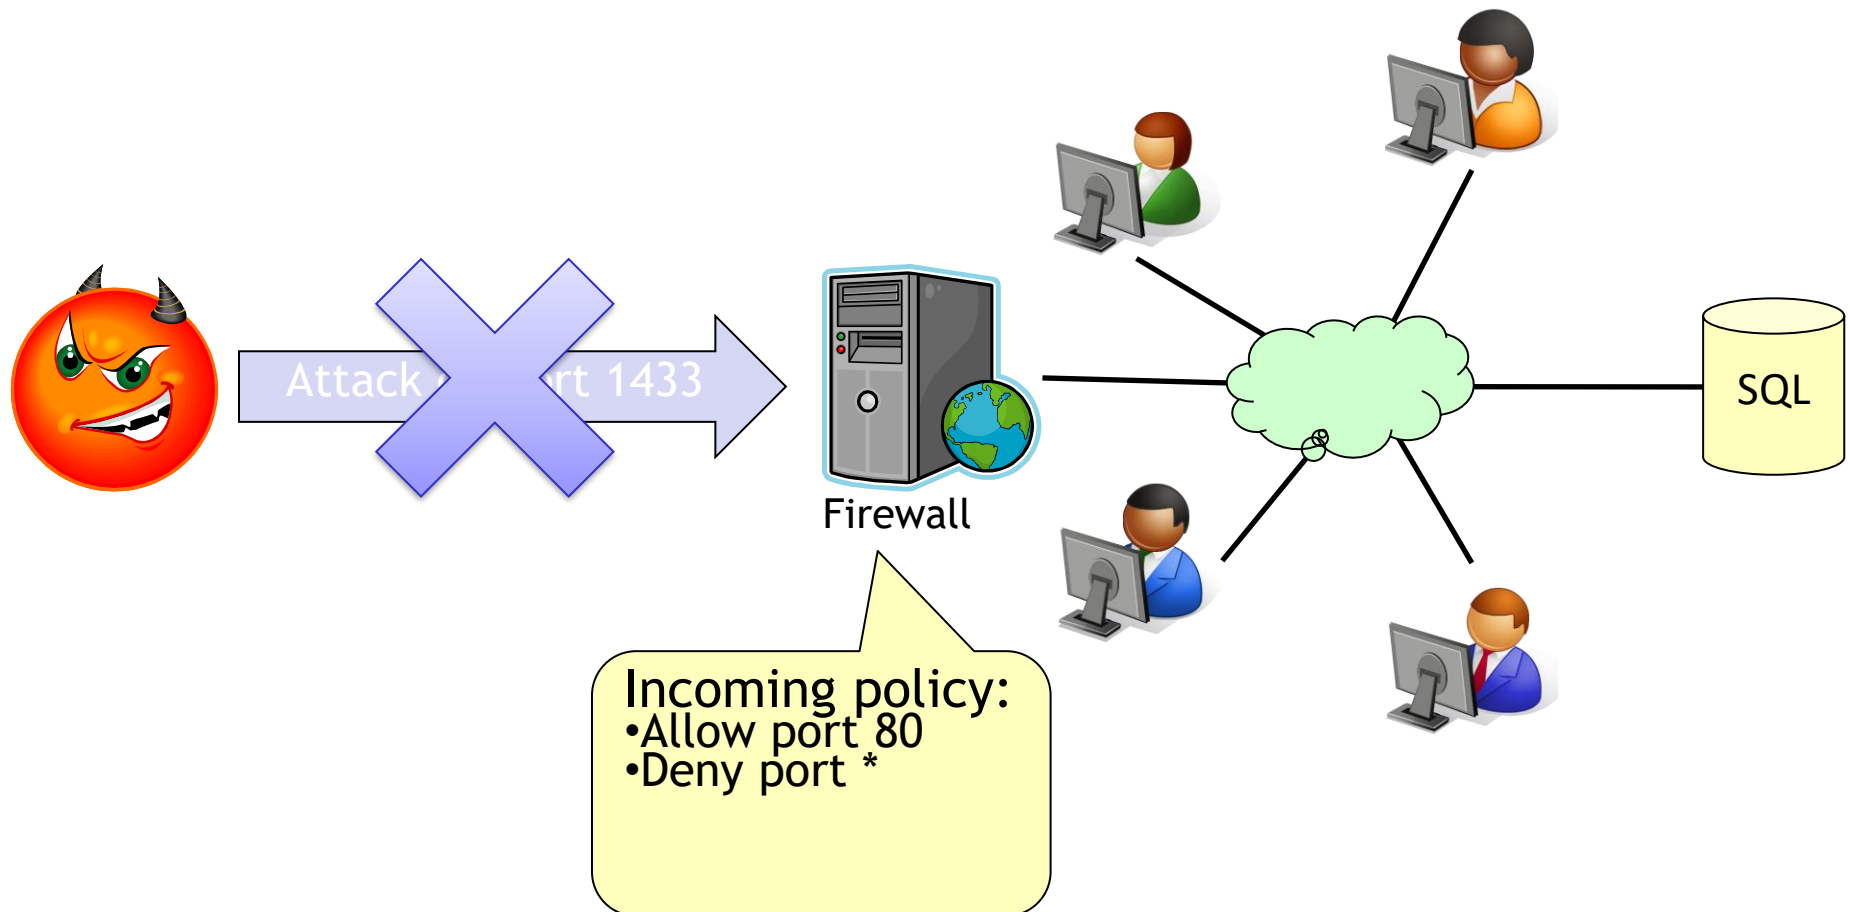
Example: Firewall protecting a small business



Principle (b): Fail-Safe Defaults

Definition: Base decisions on *permissions* rather than *exclusions*.


Example: Firewall protecting a small business



Principle (c): Complete Mediation

Definition: Every access to every object must be checked for authority.

Informally: We don't want any back doors into the system



Implications:

- Need a foolproof method for origin authentication

- Caching should be viewed skeptically (cf. DNS poisoning)

Without this, a system can offer no concrete guarantees

Principle (d): Open Design

Definition: The design of a system should *not* be secret.

Security should not be dependent on the ignorance of attackers

Attackers are often well motivated

Open design assuages skeptics and permits open review

In reality, it is impossible to keep widely-distributed software a secret

Example: DVD copy protection




Instead, use public algorithms with secret parameters

The community can verify properties of the system

Every organization can create their own secret “instance” of the system

In cryptography, this notion is called Kerckhoffs’s principle

Shannon’s Maxim: “The enemy knows the system”



Principle (e): Separation of Privilege

Definition: *Where feasible, a protection mechanism that requires two keys to unlock it is more robust and flexible than one that allows access to the presenter of only a single key.*

Examples:


Need two keys to launch a missile

Compartmentalized hulls in large ships

Witness on marriage licenses

Notarized documents

...



In computer security, the most common example of this principle is in **separation of duty**. For example, employees who create payment authorizations cannot issue checks.

Principle (f): Least Privilege

Definition: *Every program and every user of the system should operate using the least set of privileges necessary.*

This principle is **often** violated!

Windows users with “Administrator” accounts

- Accidentally deleted system files

- Installing untested programs that corrupt other user accounts

- Etc.

UNIX servers running as “root”

- Root needed to bind to port 80, but that’s it!

- Compromise of web server (public process) can lead to corruption of whole system!

Where is this principle actually respected?

- SELinux

- Attenuation of privileges in RBAC

- Restricted delegation

Principle (g): Least Common Mechanism

Definition: Minimize the amount of mechanism common to more than one user and depended on by all users.

Most common interpretation: Minimize shared channels

Memory protection


Base/bounds memory

Virtual memory

Virtual machines

VPN

...



Why?

Confidentiality/integrity protection

Information flow, side channels, etc...

Principle (h): Psychological Acceptability

Definition: *It is essential that the human interface be designed for ease of use, so that users routinely and automatically apply the protection mechanisms correctly.*

Implications:

1. If security isn't easy to use, people **won't** use it!
2. If mandatory security features aren't easy to use, people will use them incorrectly!

Example: Digitally-signed email

1. Spoofed email that could be prevented
2. Why can't Johnny encrypt?



HCISec


Implementation

An **implementation** creates a functional system based on the design.

Ideally, we want to verify two things

The design satisfies the specification

The implementation satisfies the design and (by transitivity) the specification



This is harder than it sounds!

Analyze correctness of each line of code

Preconditions

Postconditions

Proof of correctness depends on whether global preconditions to whole program hold

How do we specify these?

Are they correct?

Example: Bad compiler

Taking a testing-based approach to assurance is a popular alternative to formal verification

Intuition: If a bunch of smart people can't break into my system, I have at least some assurance that it is operating correctly

Many flavors of testing

- Regression testing/unit testing

- Red teams/penetration testing

- Fuzz testing

- ...



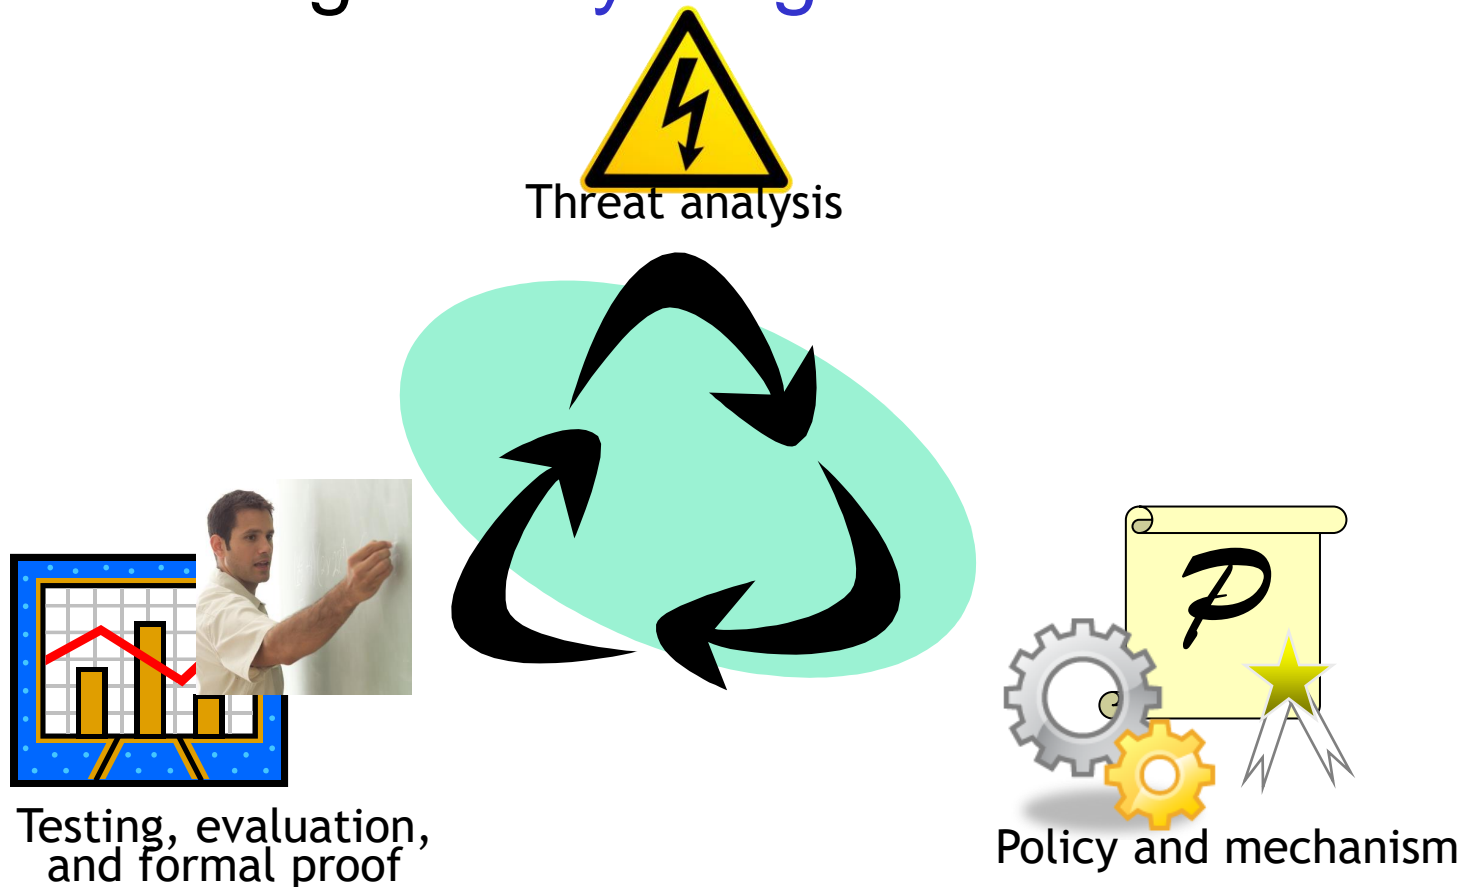
Testing **does not** provide you with a proven guarantee, but can uncover weaknesses or errors in a system

In the end, successfully resisting a **rigorous attempt** at intrusion is a good sign that things are on the right track

Systems are dynamic, so are threat models and security policies

Changing environment means changing assumptions


This can change **everything** else...



Conclusions

Computer security does not happen by accident!

Careful attention must be given to security considerations at **all** stages of the software development lifecycle



Best case scenario: Integrated process

- Organizational risk analysis and cost/benefit analysis


- Saltzer and Schroeder's design principles

- Formal verification and/or systematic testing


Our introduction is now over

Next time: Introduction to classical cryptography

Please skim over assigned reading



Project Phase 1 will be out this Friday!



Discussion

In systems that you have designed, did your team consider these types of requirements ahead of time? How can this type of security planning be incorporated into the software engineering process?

- a) Economy of Mechanism
- b) Fail-Safe Defaults
- c) Complete Mediation
- d) Open Design
- e) Separation of Privilege
- f) Least Privilege
- g) Leave Common Mechanism
- h) Psychological Acceptability