



University of  
Pittsburgh

# Applied Cryptography and Network Security

## CS 1653



Summer 2023  
Sherif Khattab  
[ksm73@pitt.edu](mailto:ksm73@pitt.edu)

(Slides are adapted from Prof. Adam Lee's CS1653 slides.)

# Announcements

- Project Phase 4 Due tonight @ 11:59 pm
  - Please schedule Project Phase 4 Demo with Pratik as soon as possible
- Project Phase 5 Due on Monday 8/7 @ 11:59 pm
  - No demo meetings for this phase
- Programming Assignment 3 canceled
  - due to delays in certificate issuance for PA 2
- Things due this Friday
  - Phase 4 Peer Evaluation Survey
  - Homework 10
  - Midterm reattempts
- OMETs Bonus
  - Entire class gets 1% bonus when response rate is 80% or more (currently at 6%)
  - OMETs due on 8/5
- Office hours updated for this week
  - Please check <https://khattab.youcanbook.me/>
  - You can still request (by email) office hours outside those times

# Final Exam

- Take home: download from GradeScope, (print), solve, (scan), and upload to GradeScope
  - open-book and open-notes
  - some overlap with midterm topics
- Please check study guide on Canvas
  - Review video will be posted soon
- Exam will be posted tomorrow
  - You will have until Sunday 11:59 pm to finish and upload the exam
- No class on Wednesday

# Private Routing

(Very quick) overview of routing

**Discussion:** What are desirable security properties in routing?

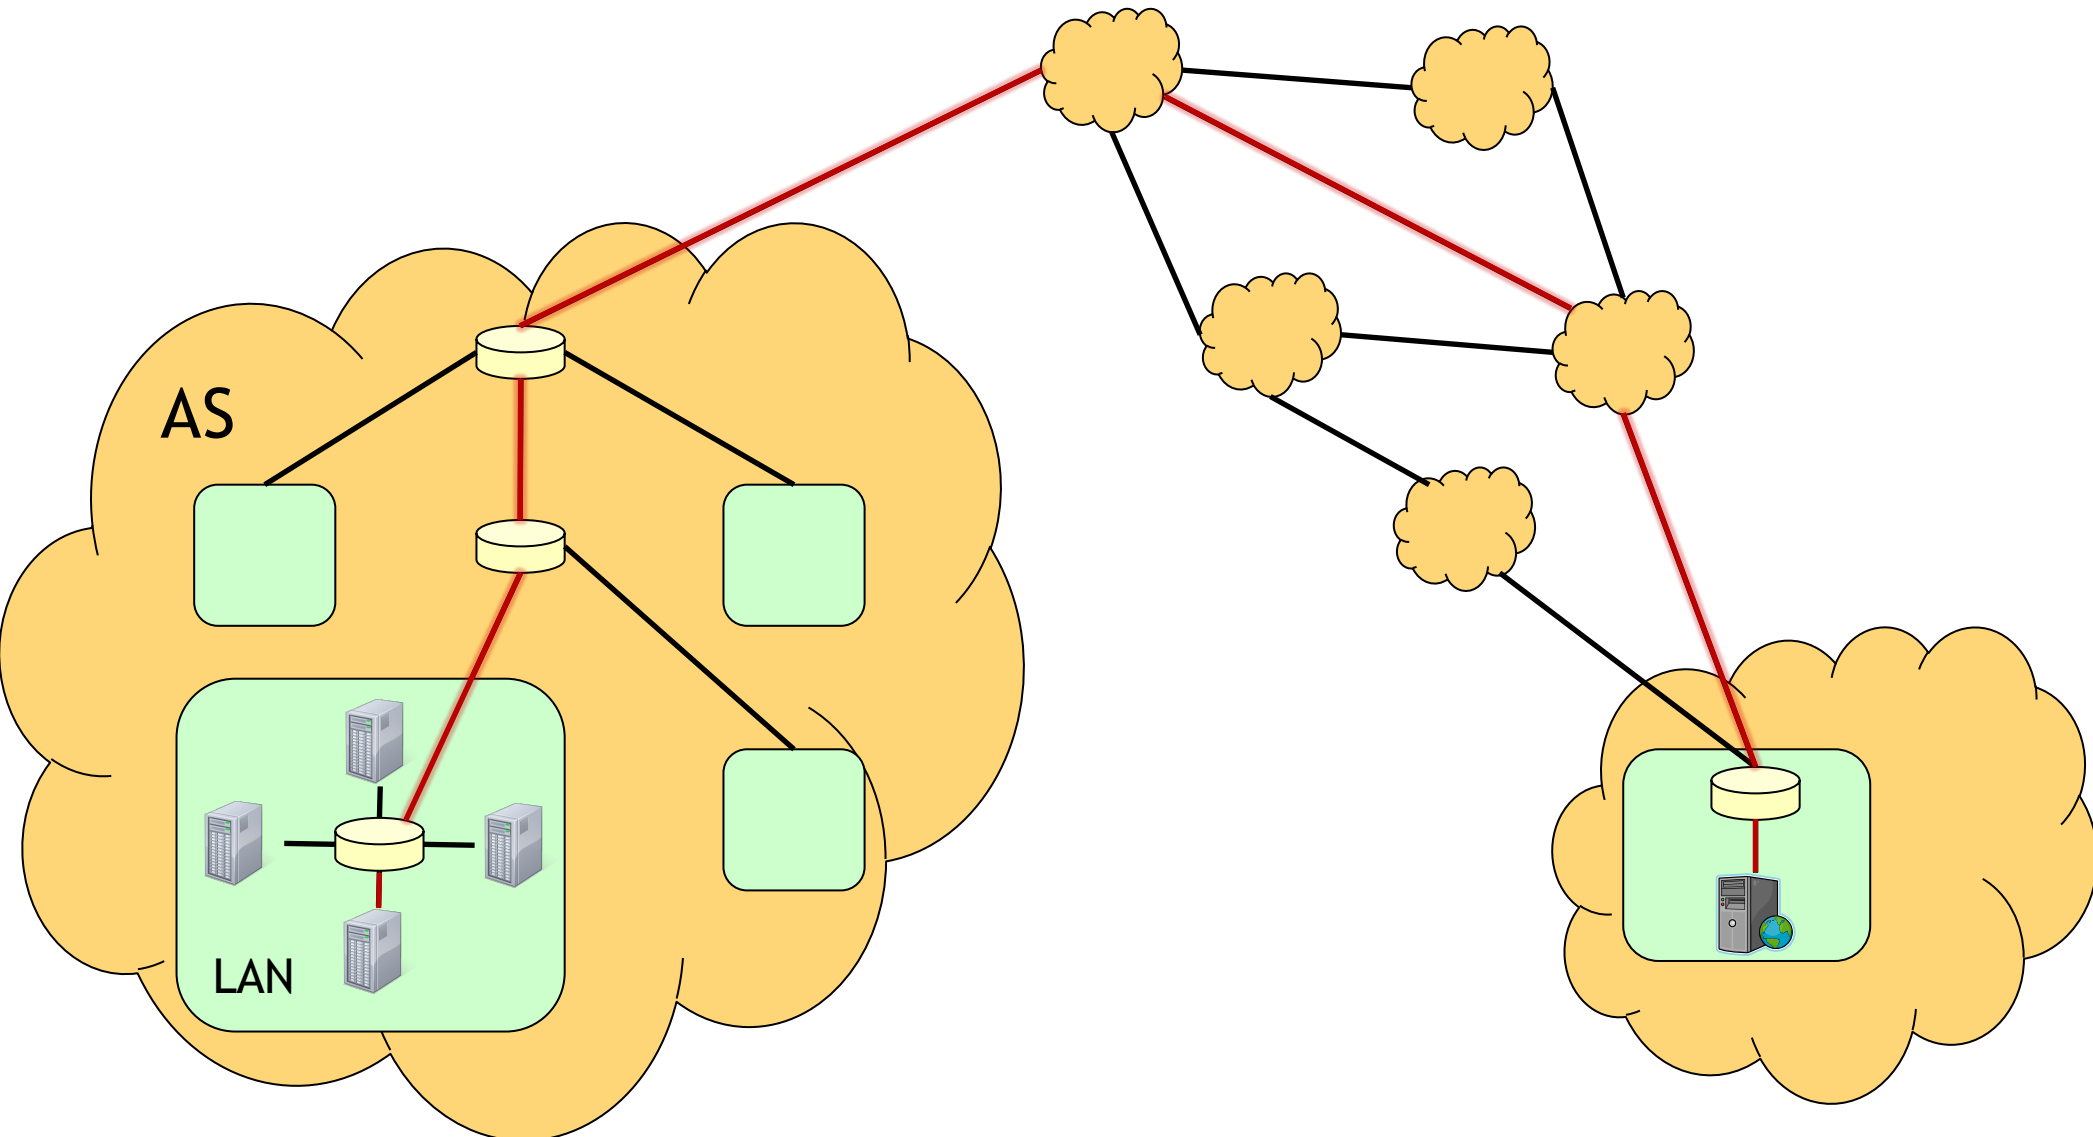
Confidentiality & **Firewalls**

Integrity & **DNSSEC**

Privacy & **Crowds, Tor**

# Network overview and routing

The Internet is a network of networks



# Networked systems represent a change in paradigm from the “old days”

---

*What makes networked systems vulnerable to attack?*

---

## 1. All traffic is routed in public

- **Early days:** Principals largely trusted, plaintext traffic → IP simplified
- The result is that communicating on the Internet is like shouting in a crowded room, but users think they're sending private messages

## 2. Traffic flows across many administrative domains

- High degree of exposure across potentially untrusted territory

## 3. There may be multiple routes between two points

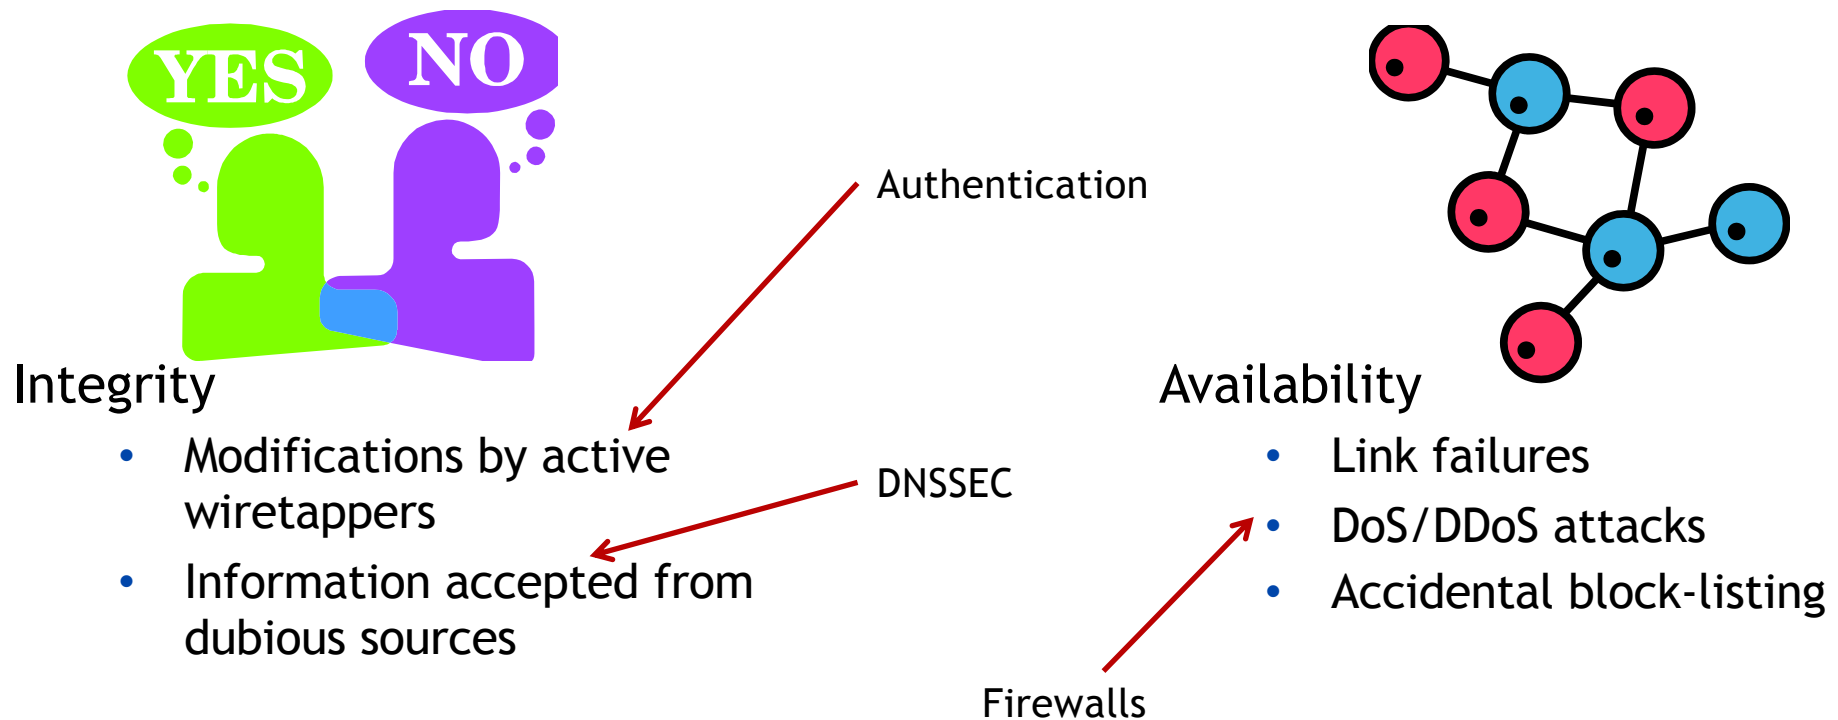
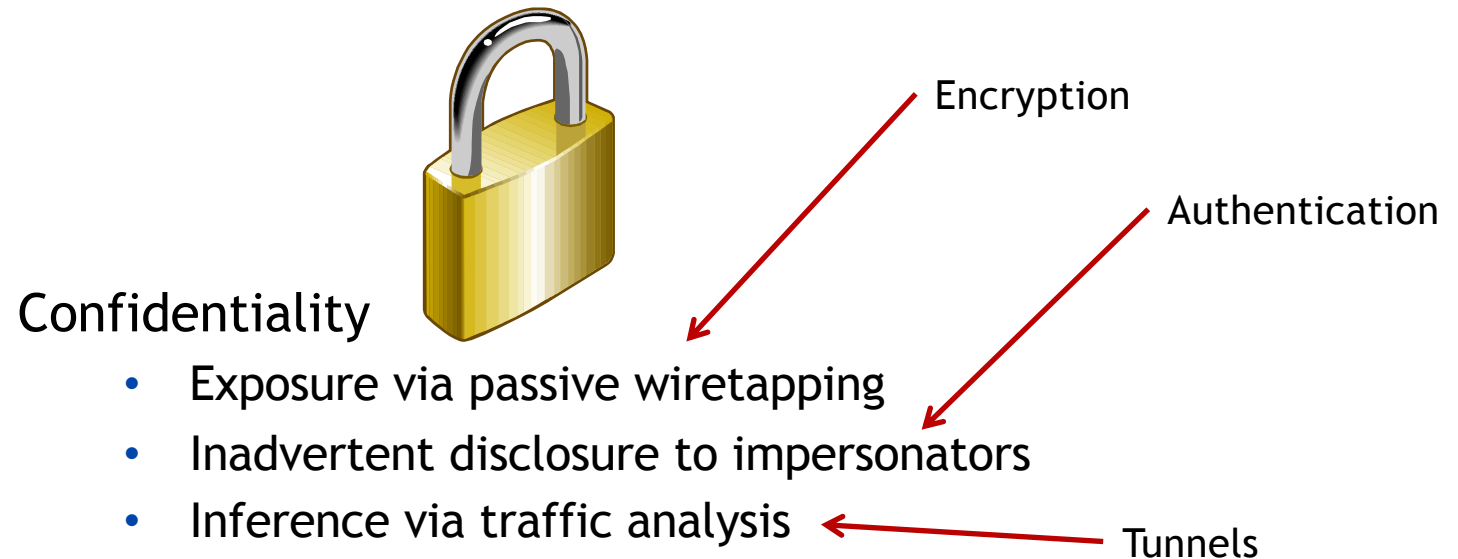
## 4. Principals have some degree of anonymity in the system

- Users don't need physical access to resources
- Simplified IP means that addresses are not bound to identities
- Worse yet, **stepping-stone** attacks!

# Discussion Question

*Given these vulnerabilities, what are some threats against network systems?*

# A sampling of threats for networked systems...

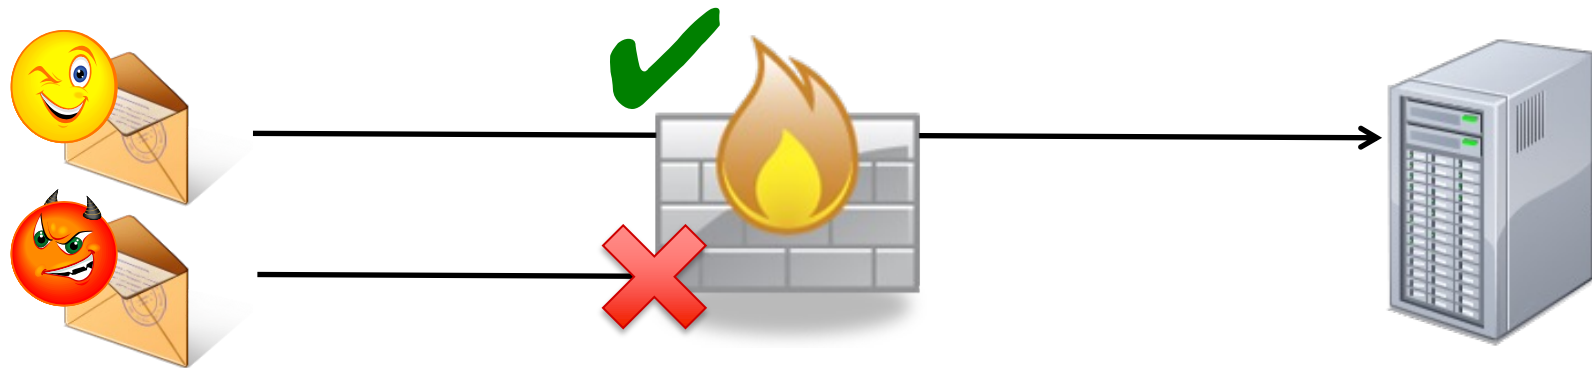




# **FIREWALLS**

# Firewalls are *the* ubiquitous network security mechanism

First proposed in 1988, firewalls act like networked reference monitors



Why are firewalls necessary?

- The early days were a more optimistic time!
- Allowing all traffic to be routed between all points greatly simplified the design of TCP/IP

Several types of firewalls are in widespread use

- “1G”: Packet filtering firewalls
- “2G”: Stateful packet inspection firewalls
- “3G”: Application-layer firewalls
- Machine-learning powered firewalls

# Packet filtering firewalls are the simplest type of firewall

As packets arrive, a pass/drop decision is made by inspecting headers:

- IP protocol (e.g., TCP, UDP, ICMP, ...)
- Source and destination IP addresses
- Source and destination ports (if applicable)

Most firewalls allow rules to either **grant** or **deny** access

- allow TCP \* \* 123.4.5.6 80
- deny \* 67.8.9.10 \* \* \*

Everyone should have access to the web server

This might be a compromised server attempting to infect us with a worm

Packet filtering firewalls usually run on dedicated devices with very simple operating systems

- No linkable libraries, compilers, etc.
- Minimal system tools
- Sometimes, no system accounts!
- Why? Minimal TCB

If compromised, make it hard for the attacker to stage more attacks

Make it hard to change the system configuration

No accounts means no privilege escalation!

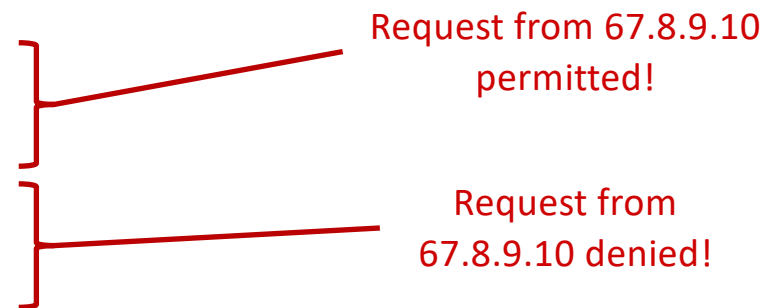
# Simplicity gives rise to complexity

Packet filters often have enormous numbers of rules

- In large corporations, firewalls with 10k rules are not uncommon
- Anecdotally, some firewalls have up to 50k rules!

Worse, often the order of rules matters

- allow TCP \* \* 123.4.5.6 80  
deny TCP 67.8.9.10 \* 123.4.5.6 80
- deny TCP 67.8.9.10 \* 123.4.5.6 80  
allow TCP \* \* 123.4.5.6 80



What happens if a packet does not match any defined rule?

- Security administrators like **default deny** policies
- Users like **default permit** policies

This complexity has sparked many interesting research directions

- What high level policy is encoded by a set of firewall rules?
- What are the effects of changes to firewall policies?
- How can rules be more efficiently organized?
- What is the net effect of multiple layered firewalls?

# More complex firewalls allow more expressive policies, at a cost

“2G” **Stateful firewalls** keep track of connection state

- Distinguishes new connections from existing connections
- Example: FTP
  - Commands over port 21, transfers over arbitrary high ports
  - Consider single packets → **deny**
  - Associate with existing connection → **allow**
- Vulnerable to DoS via filling connection state memory

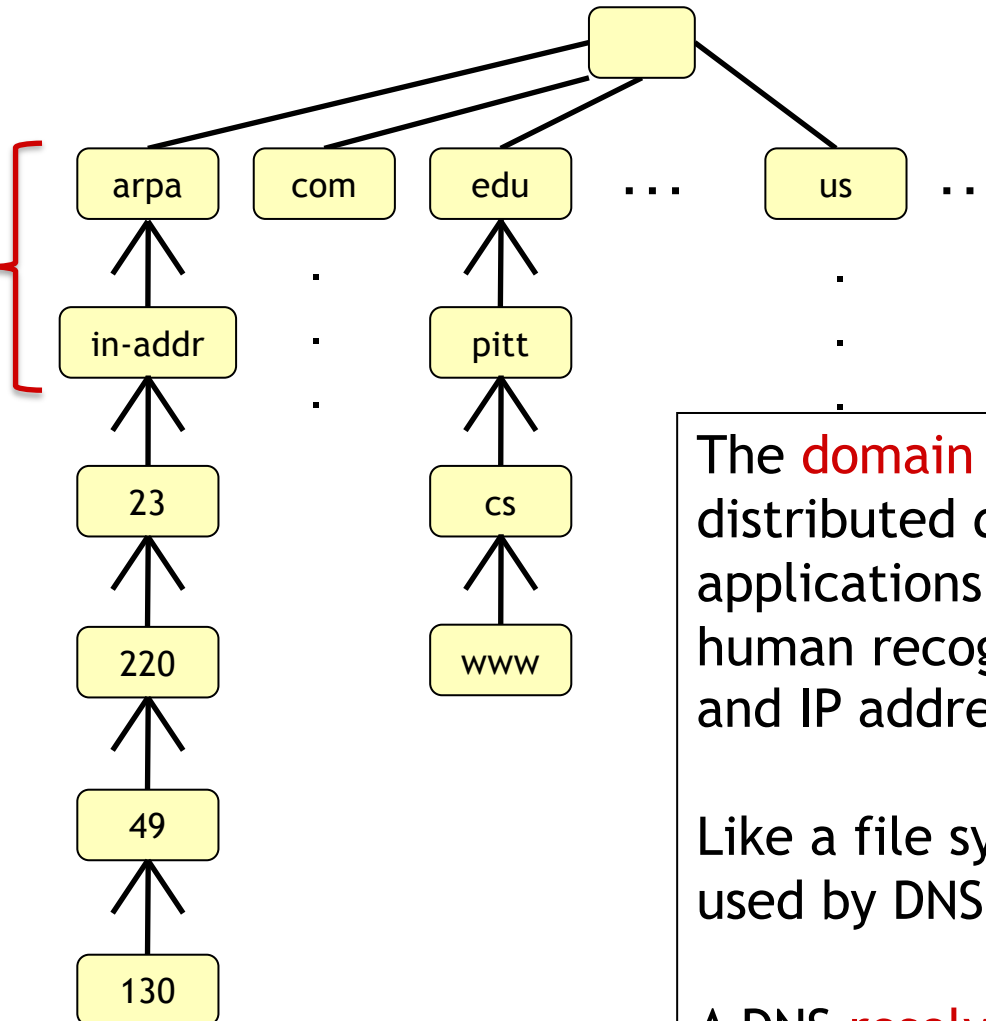
“3G” **Application-layer firewalls** inspect application data in packets

- Ability to interpret application layer allows tighter control
- Detect forbidden protocols over allowed ports, misuse of allowed protocols
- Bind application-layer identity (e.g., FTP username) to originating IP
- Scan for viruses and worms
- Much higher overhead than lower-layer firewalls

# **DNS SECURITY**

# DNS is the mechanism through which host names are resolved into IP addresses

\*.in-addr.arpa is used for reverse lookups

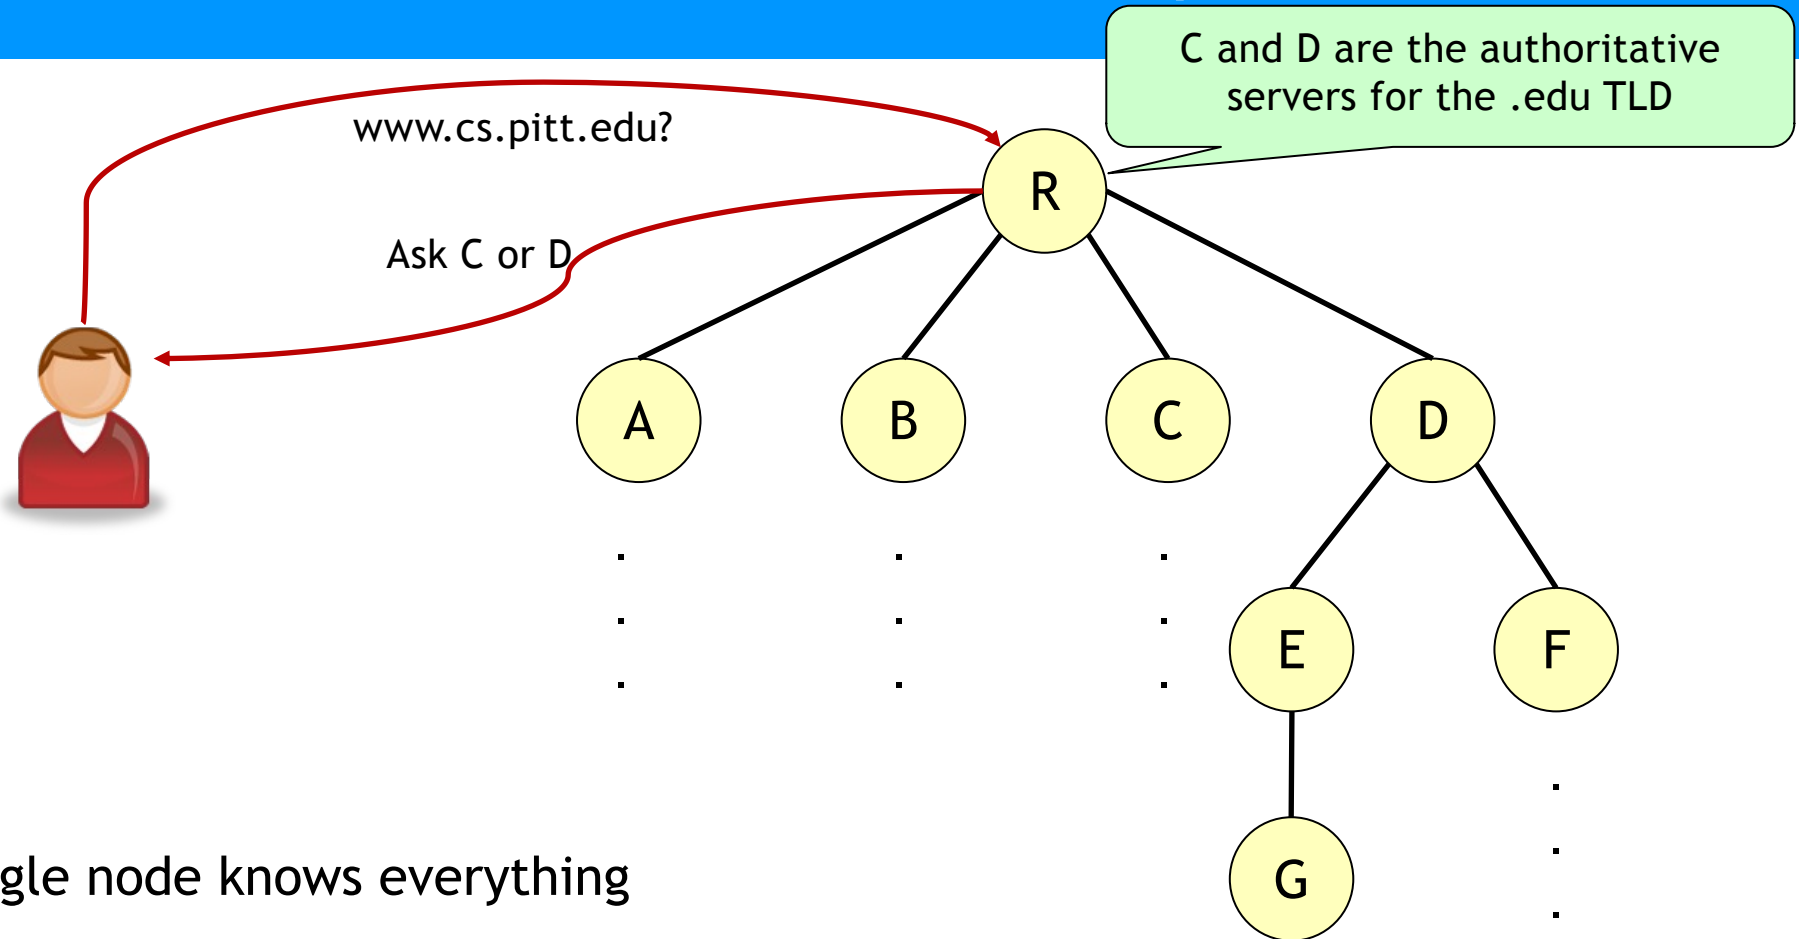


The **domain name system** (DNS) is a distributed database that allows applications to map between human recognizable domain names and IP addresses

Like a file system, the namespace used by DNS is **hierarchical**

A DNS **resolver** can start at a known “root” node and issue recursive queries until it finds the address that it is interested in

# DNS Resolution Example

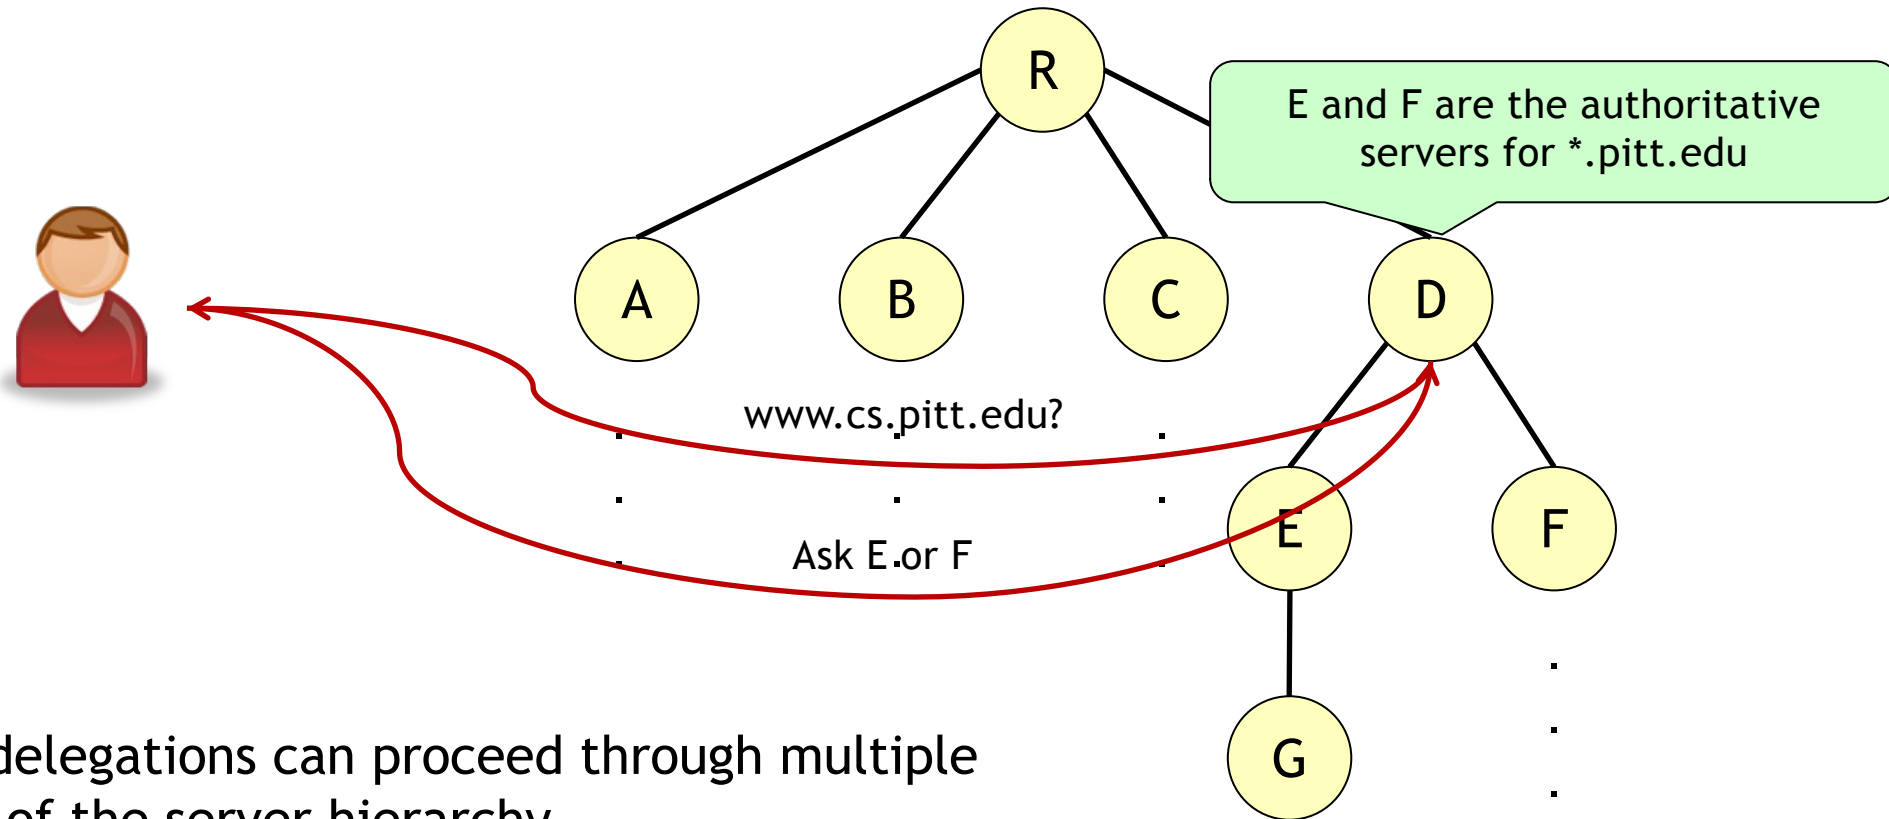


No single node knows everything

**NS records** identify the authoritative servers for a particular zone

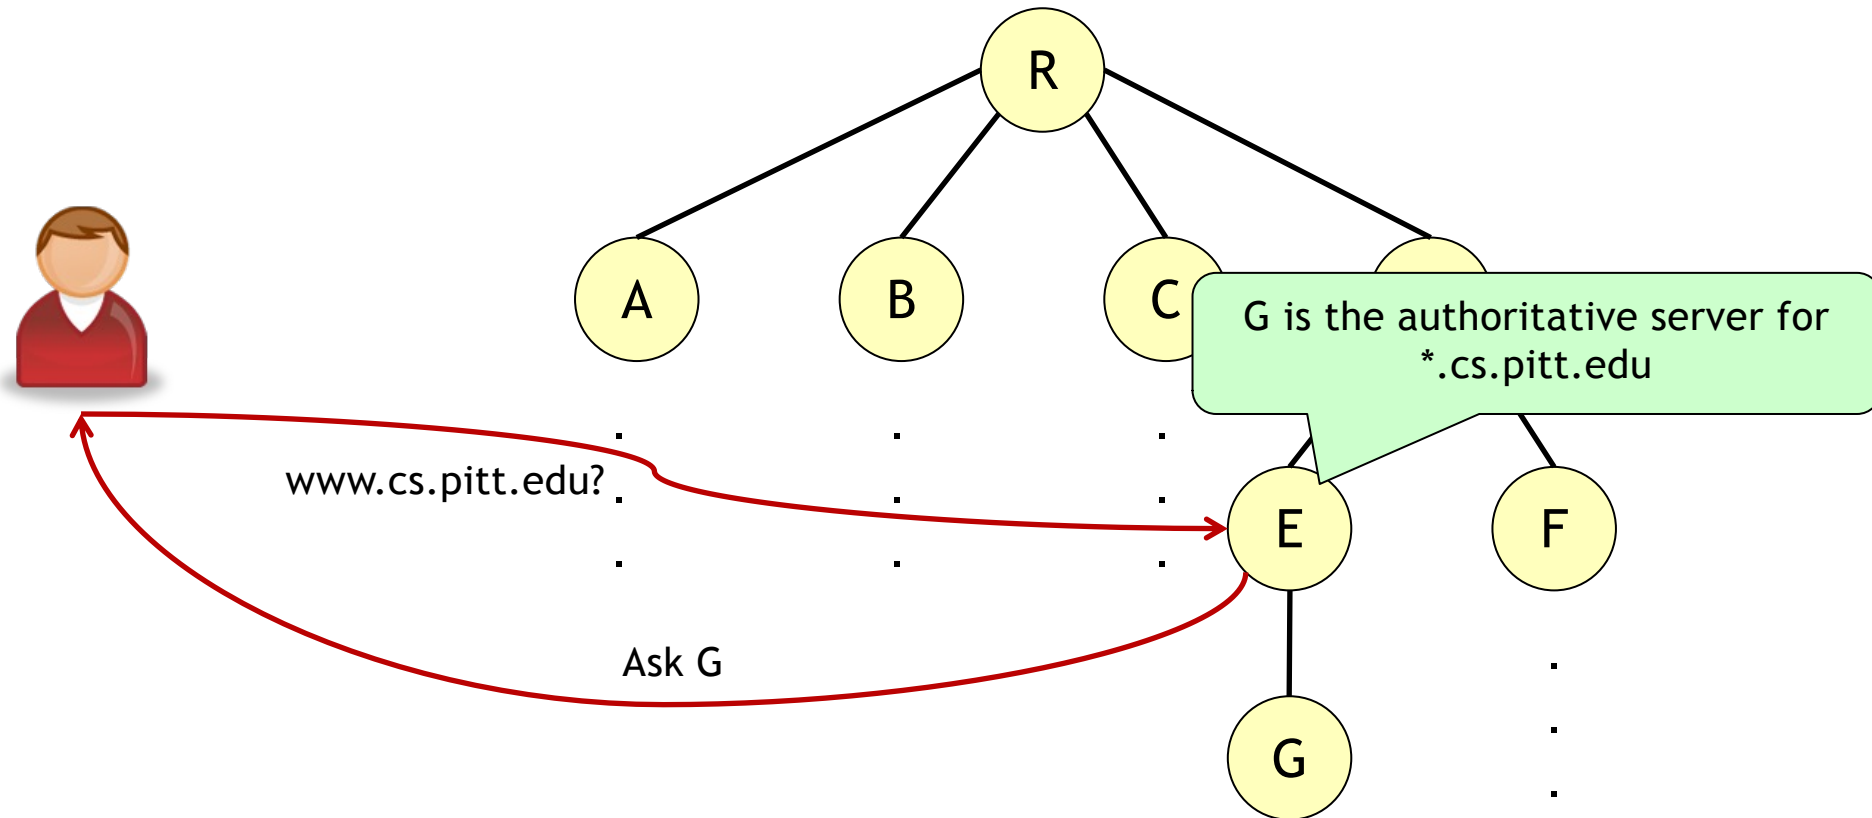


# DNS Resolution Example

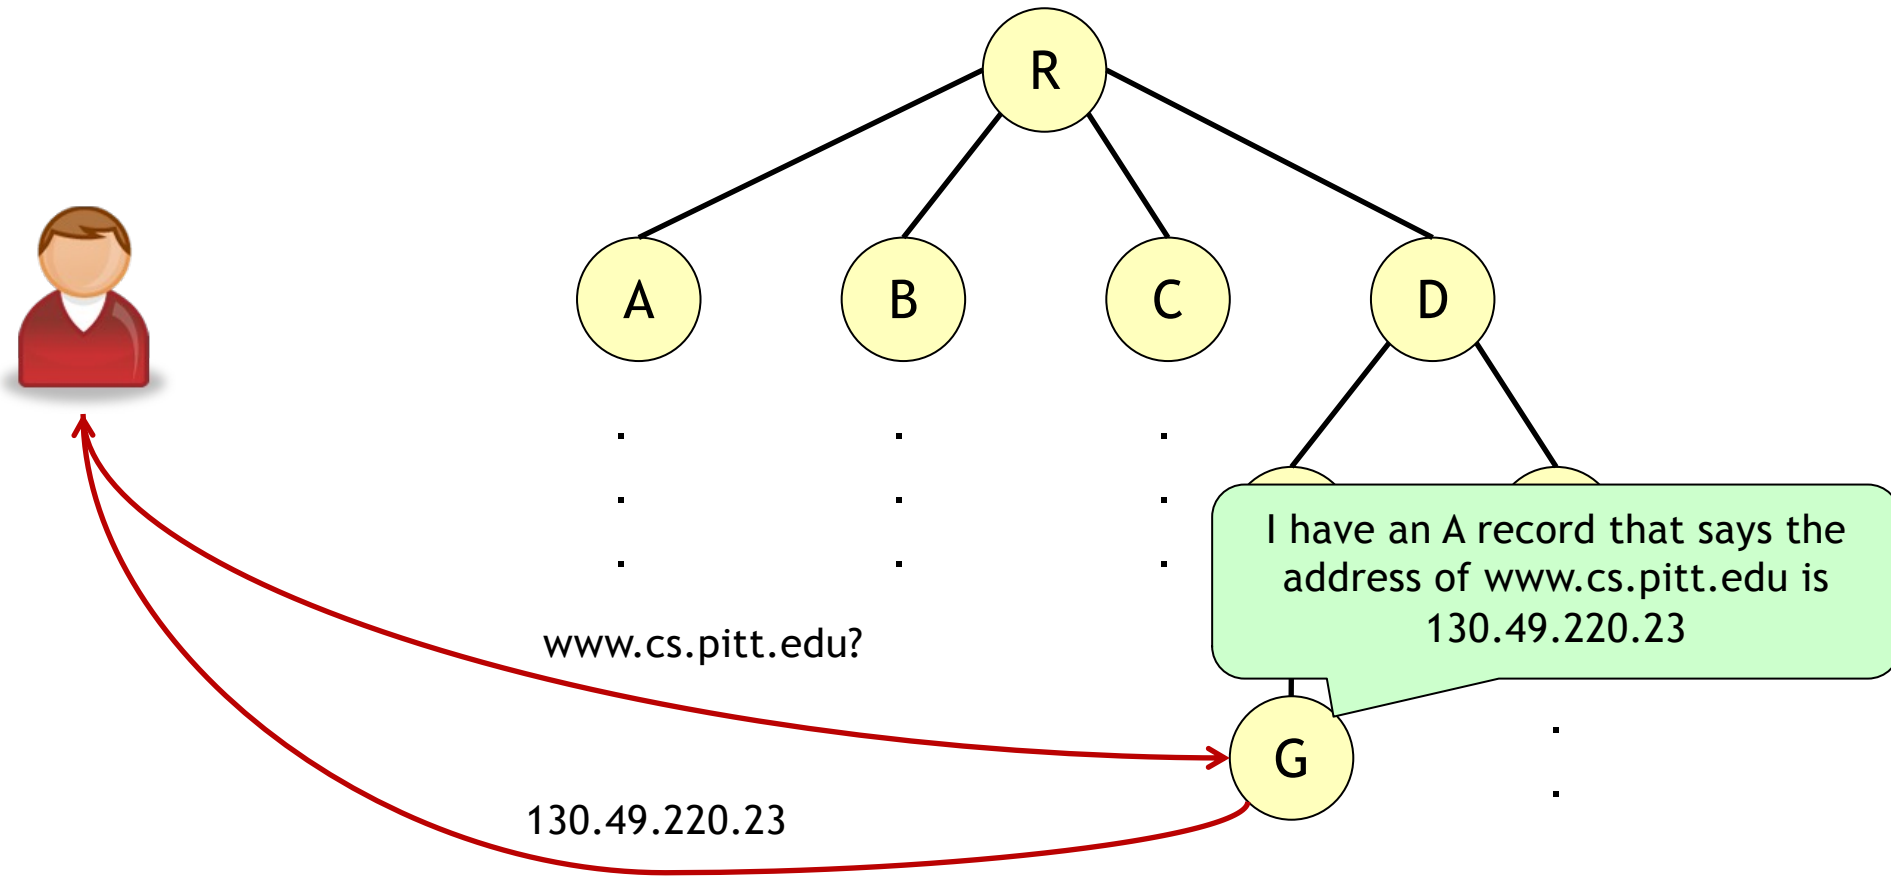


Zone delegations can proceed through multiple levels of the server hierarchy

# DNS Resolution Example



# DNS Resolution Example

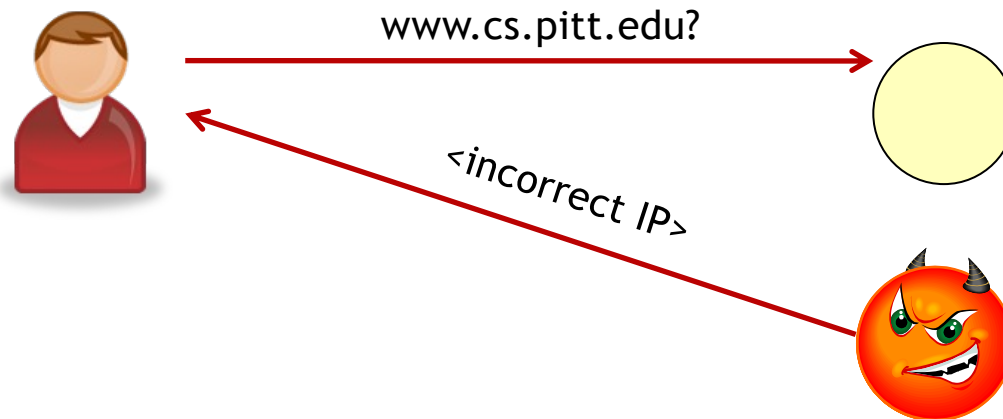


**A records** are used to manage the bindings between a domain name and an IP address

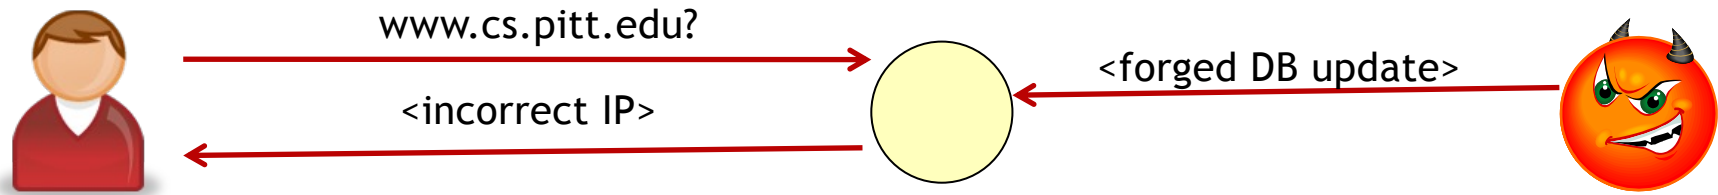
# How could such a simple system have problems?

DNS **lacks a strong authentication mechanism** that can be used to verify the integrity of records returned

*Example:* Forgeries injected at the last hop



*Example:* Incorrect zone transfers



# DNSSEC has been proposed as a solution to these types of problems

**Insight:** DNS servers form a hierarchy, PKIs also form a hierarchy!

To hybridize DNS and PKI, new DNS resource records were introduced

- **DNSKEY** records are used for each node to store its public key
- **DS** records are used to delegate authority to other nodes
  - Kind of like an NS record
  - A signed binding between a node IP, zone, and public key
- Furthermore, **all** DNS resource records can be **signed**

To use DNSSEC, nodes need the keys of trusted root DNS servers

The result is authenticated responses!

- $\text{DNSKEY} \rightarrow [\text{DS} \rightarrow \text{DNSKEY}]^* \rightarrow \text{Signed RRset}$

Trust root

Signed A records

Zero or more **signed** delegations  
through the zone hierarchy  
(like PKI certificates)

# So, why haven't we seen widespread deployment of DNSSEC yet?

***Reason 1:*** Some people want to see the protocol more thoroughly vetted

- Earlier versions of DNSSEC had problems
- Complex 6-message protocol to update node keys
- Every record on the child server had to be signed by parent
- **Result:** Couldn't scale to Internet sizes

***Reason 2:*** US representation on the Internet

- Many root servers are located in the US
- Other nations perhaps uneasy about giving so much control over cryptographic keys to one political entity

***Reason 3:*** Privacy

- DNS zone data is usually kept private (phone book vs. reverse lookup)
- DNSSEC also needs signed failure messages
- Combination of signed failures and signed successes allows adversaries to look up **all** data regarding a zone

# **PRIVATE ROUTING**

# Application-level encryption cannot fully protect privacy

## Open System Interconnection Reference Model

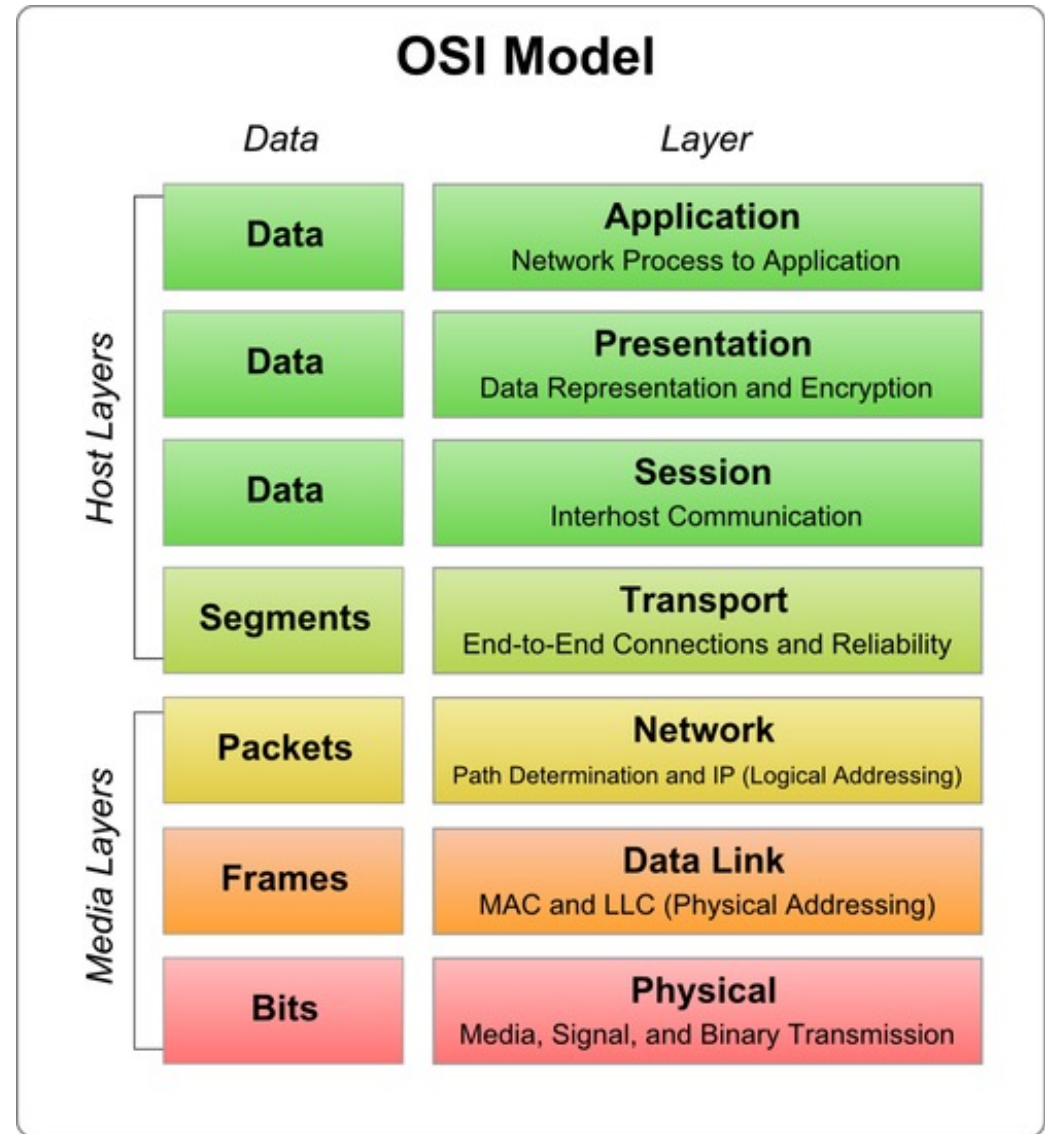
Host layers can easily be encrypted

- e.g., SSH, TLS

What about info at the network layer?

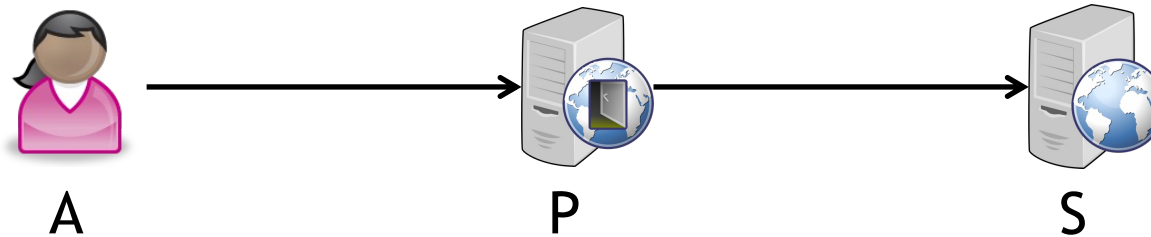
- e.g., IP addresses?

Passive adversary can infer a lot from this information!





# A single trusted proxy server can mitigate some of these problems



By using a proxy, Alice obscures the fact that she is communicating with S

Passive eavesdropper (e.g., ISP) never sees a packet identifying both Alice and S

S sees incoming traffic from P, not Alice

**Question:** What does the proxy see?

# Crowds can reduce the amount of trust placed in a single proxy

**Intuitively:** Hide in a *crowd* of other users

Two types of entities

- **Jondos:** clients wishing to hide in the crowd
- **Blenders:** trackers maintaining a list of active jondos

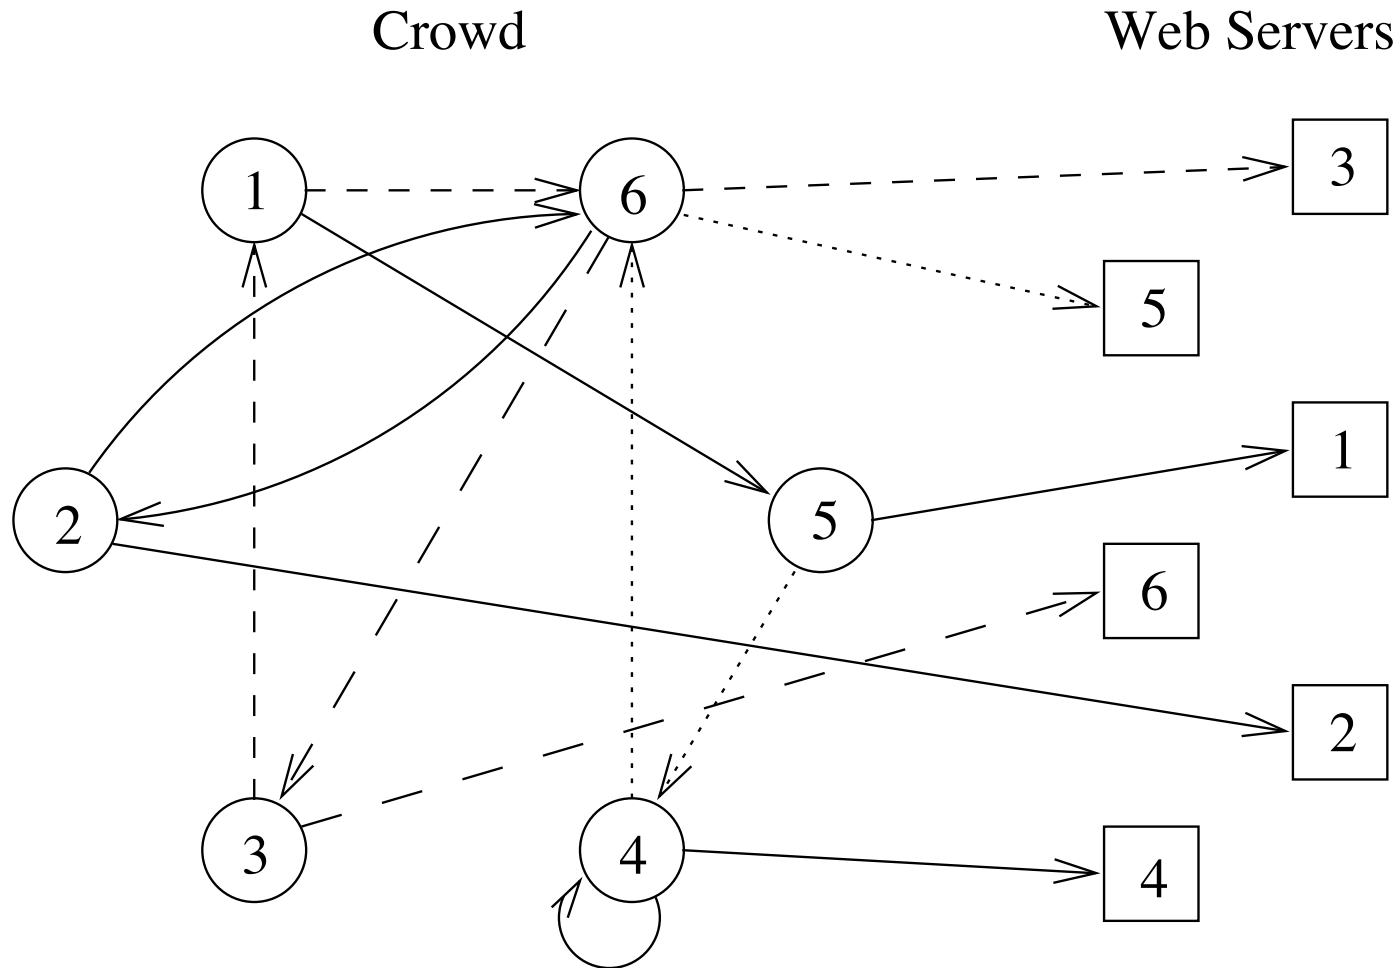
New Jondo registration with a blender

- Blender assigns new jondo a symmetric key
- Blender notifies other jondos and distributes key

Path construction when sending messages

- Each message is sent via a random path through the crowd
- When a jondo receives a message, it either:
  - Forwards it to another jondo (probability  $p_f > \frac{1}{2}$ )
  - Sends it to its destination (probability  $1 - p_f$ )

# Paths through the crowd



# How much privacy can crowds give Alice?

Passive eavesdropper:

- Sees an outgoing packet to another jondo
- Knows a packet originated with Alice if there is no corresponding incoming packet
- With  $1 - p_f$  chance, sees Alice submit message to endpoint

Endpoint:

- Sees an incoming packet, could be from anyone in the crowd

Other jondos:

- Cannot tell whether packet originated with Alice or if she is passing it along
- Colluding jondos? Alice is safe if:

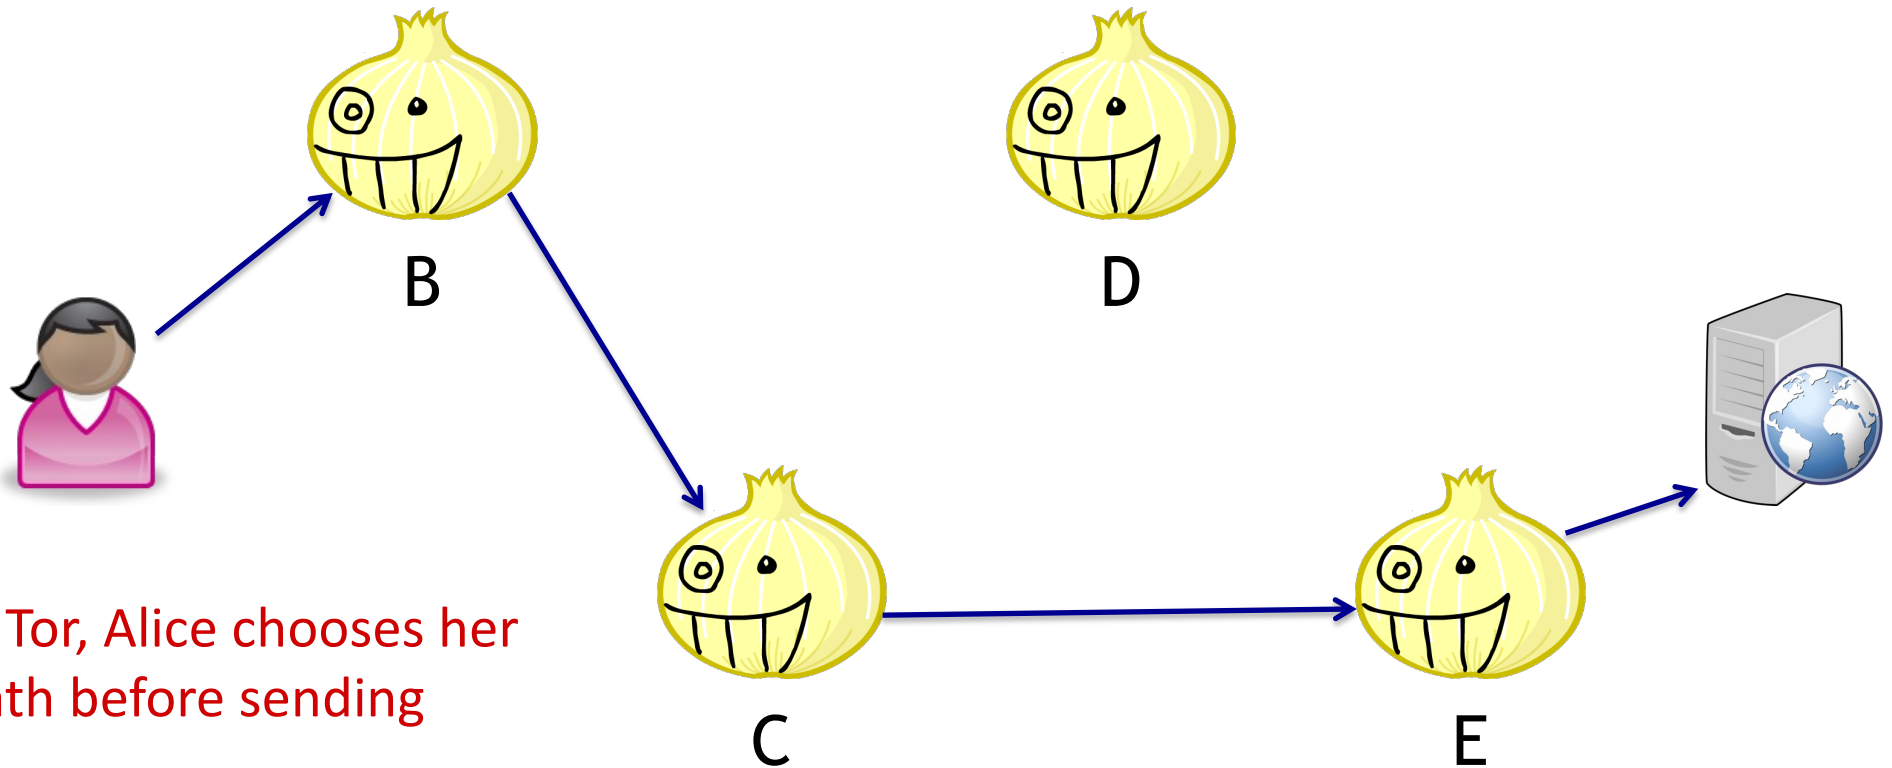
Number of jondos in crowd  $\rightarrow n \geq \frac{p_f}{p_f - \frac{1}{2}} (c + 1)$

Probability of forwarding  $\rightarrow p_f$

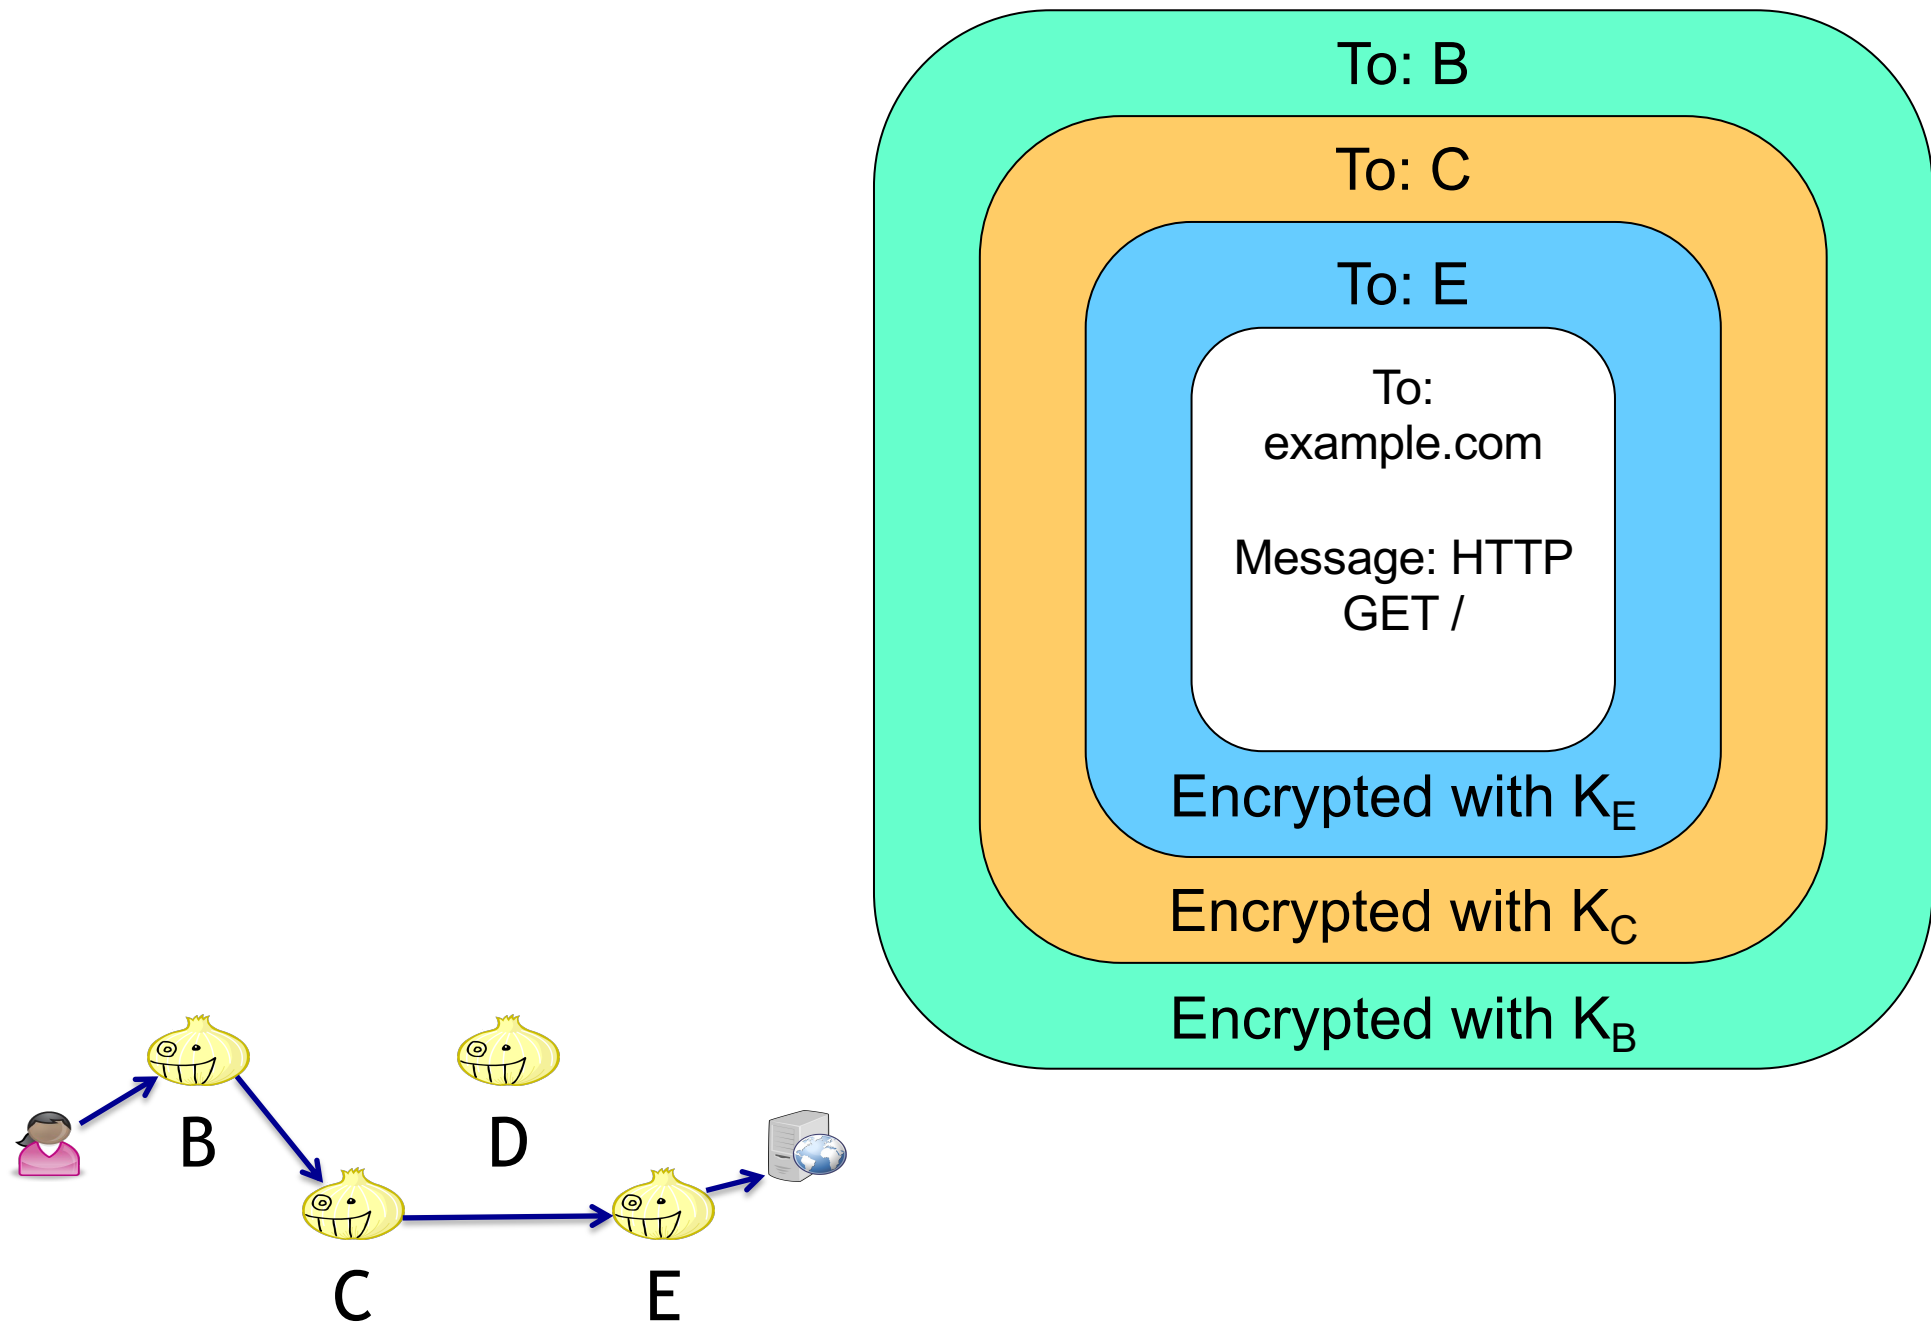
Number of colluding jondos  $\rightarrow c$

# Can we introduce encryption to increase privacy guarantees?

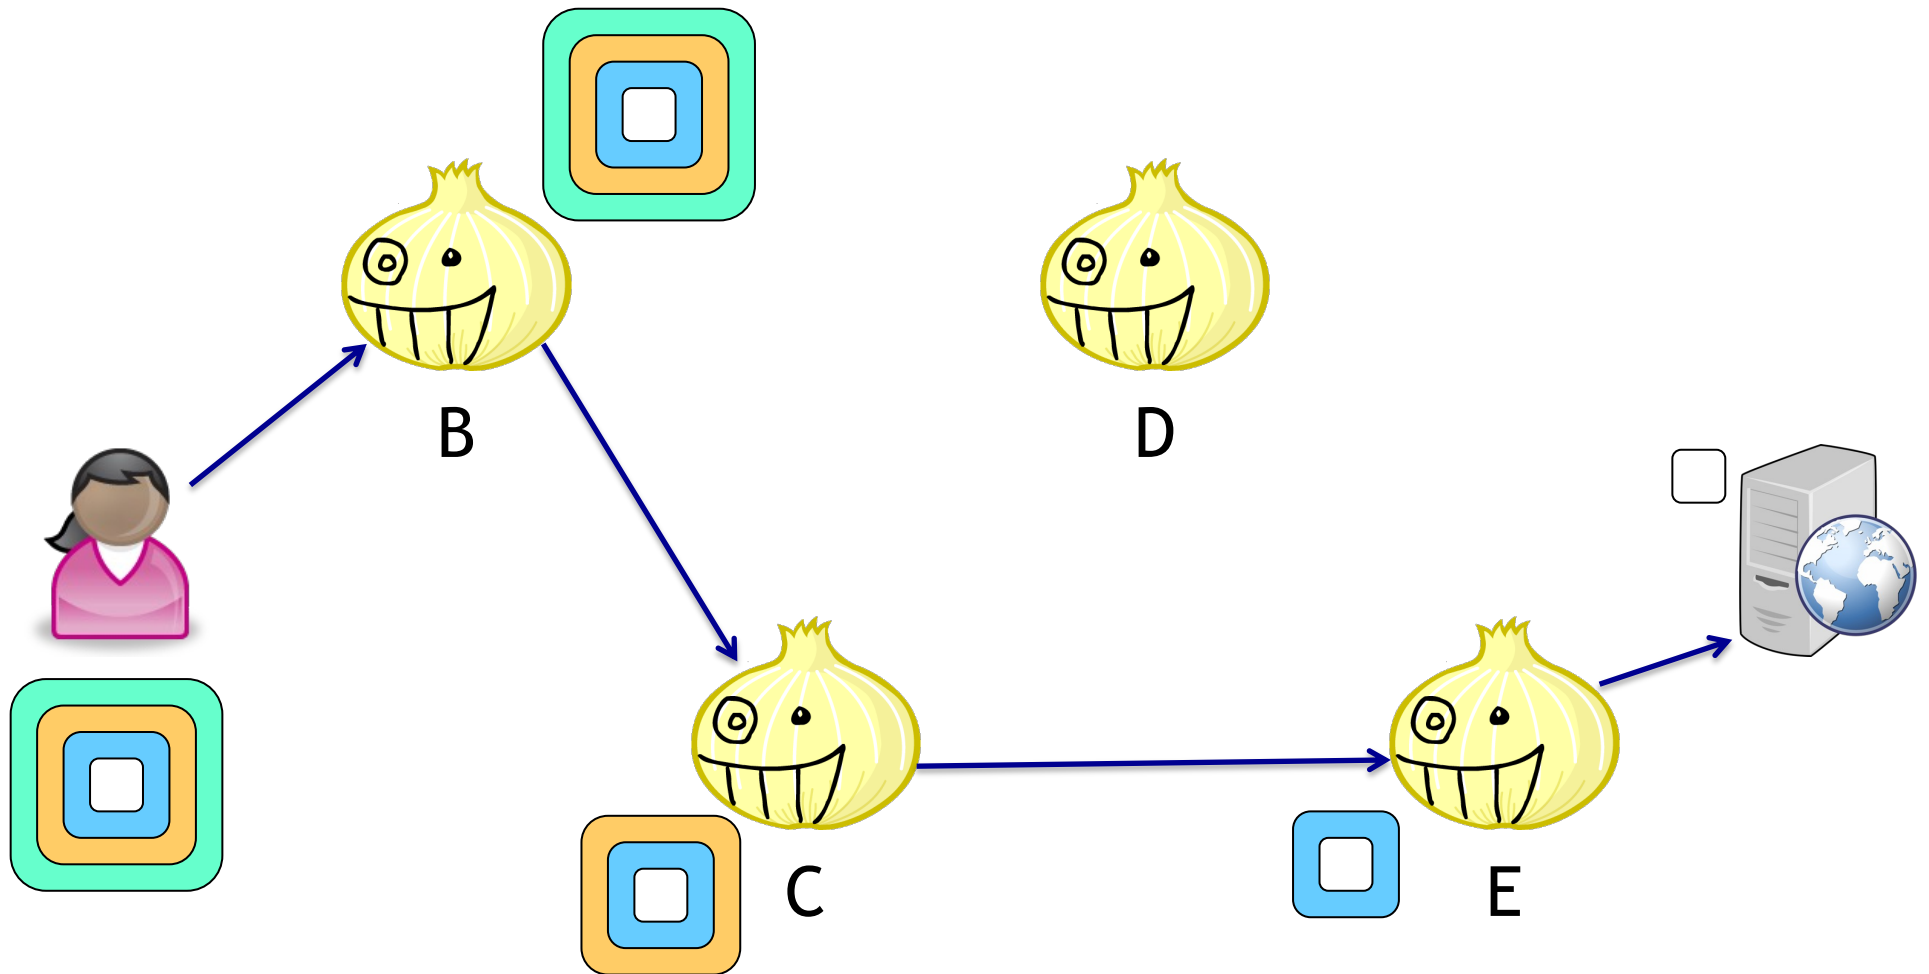
This is the idea behind onion routing (**Tor**)!



A message in Tor is protected by multiple layers of encryption, starting at the end of the path



# How does Alice send an onion message?



# How much privacy can Tor give Alice?

## Passive eavesdropper:

- Sees an onion packet outbound to an onion router
- Might know the packet originated with Alice, if Alice is not an onion router

## Endpoint:

- Sees an incoming packet from the Tor exit node

## Onion routers:

- First router sees that the message is from Alice (and the next onion router)
- Exit node sees that the message is intended for the endpoint (and the onion router that sent it)
- In between, onion routers see nothing except previous/next routers



# What threats still exist against Tor?

## Global passive adversary

- Can observe all traffic between all nodes
- By observing timing of packets, might be able to statistically determine who is talking to whom
- **Question:** How can this be prevented?

## Control of exit node

- Exit node sees plaintext message and destination
- What if this packet contains any identifying information?
- Intuition behind **Bad apple attack**: Application sends real IP in onion message, then exit node can identify the client

## Onion router collusion

- What if an attacker controls many onion routers?

# Summary

Network security is a huge topic; one size rarely fits all

**Firewalls** are used to repel intruders

Different classes of firewalls trade off efficiency for policy expressiveness

**DNSSEC** is a proposal to maintain lookup integrity

However, concern over its implications for privacy and political power have inhibited its wide-scale deployment

**Crowds** and **onion routing** protect message flow privacy

Privacy is not a boolean property—it is more like a continuum

# Data Privacy

What is data privacy? Why should I care?

Models for data privacy

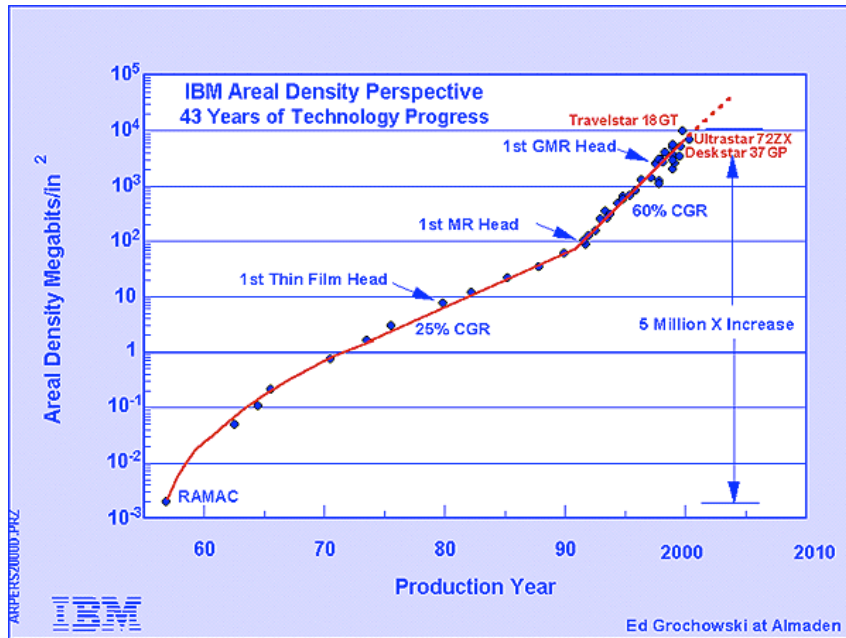
- Anonymize and release
- Mediated query processing
- Outsourced data management

*Case study:* k-Anonymity

- How does it work?
- Why **doesn't** it work?

Future directions

# Data, data everywhere!



Hard drive sizes are absurd!

- Capacity increasing
- Cost decreasing
- **Example:** 1TB backup drive costs < \$10

**Thought:** Why delete anything? Just as easy to keep it all...

**Result:** Our whole lives are on disk!

These days, “data” means more than “documents”

- Electronic health records
- Pay as you go car insurance
- Browsing/shopping histories
- Location-based services
- Social networking blunders
- ...



**Result:** Compromise can hurt more than productivity

# We can learn a lot from this data



## Google: Advertising and search

- Why are Google's services free?
- Because they use your information to intelligently place ads!
- Portions of this data is also available to you (cf. Analytics)

## Walmart: Marketing experts

- Over 580 TB in 2006, hosted on 1000 processor system
- Data used to predict/control inventory, coordinate with suppliers, and adjust to local trends



## Medical data and imaging

- Medical data mining
- [Google flu trends](#) (stopped due to privacy concerns)
- Drug and prosthesis design
- ...



# Widespread data availability is not always a good thing, though...



August 2006: AOL releases search data

- 20,000,000 search keywords
- Over 650,000 users
- 3 months worth of records

**Intended use:** Learning about search patterns

**Result:** Records for individual users were recovered!

---

October 2006: Netflix releases movie rating data

- 100,480,507 ratings that 480,189 users gave to 17,770 movies
- $\langle \text{User, Movie, Date, Rating} \rangle$  tuples

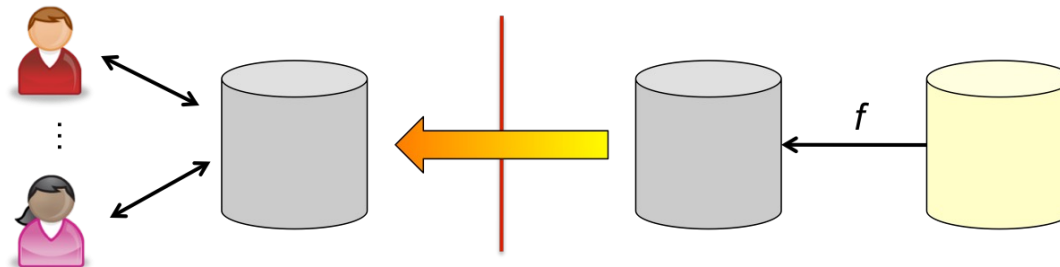
**Intended use:** Developing and testing new collaborative filtering algorithms



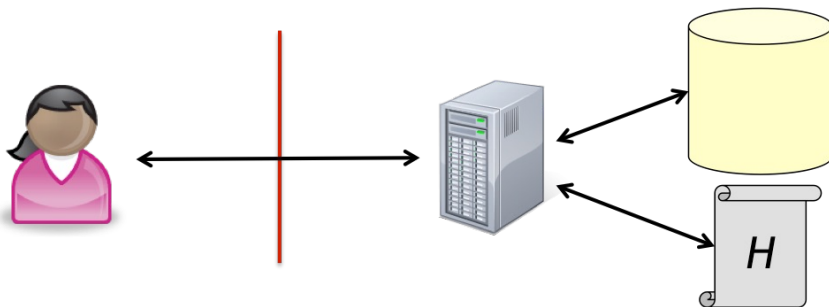
**Result:** Records for individual users were recovered!

# There is a need to balance privacy and availability when releasing data

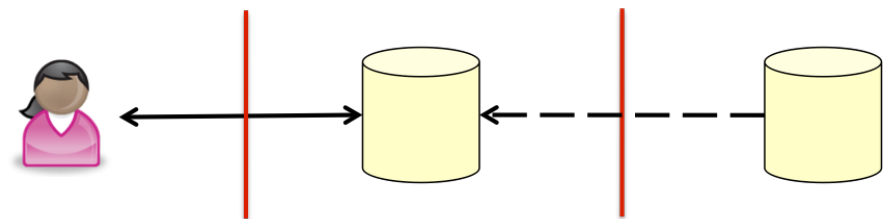
*Today, we'll talk about three privacy models for data*



**Anonymize and Release**

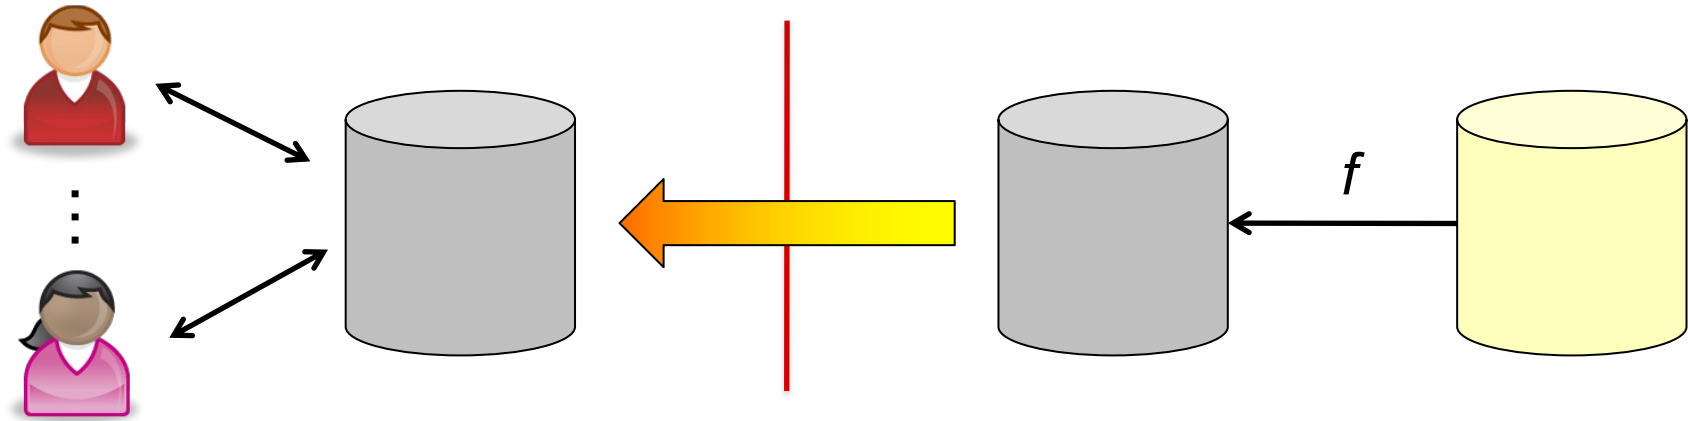


**Mediated Query Processing**



**Outsourced Data Hosting**

# Anonymize and Release



Rather than releasing the original dataset, data providers release a modified version of the dataset to the public/analysts

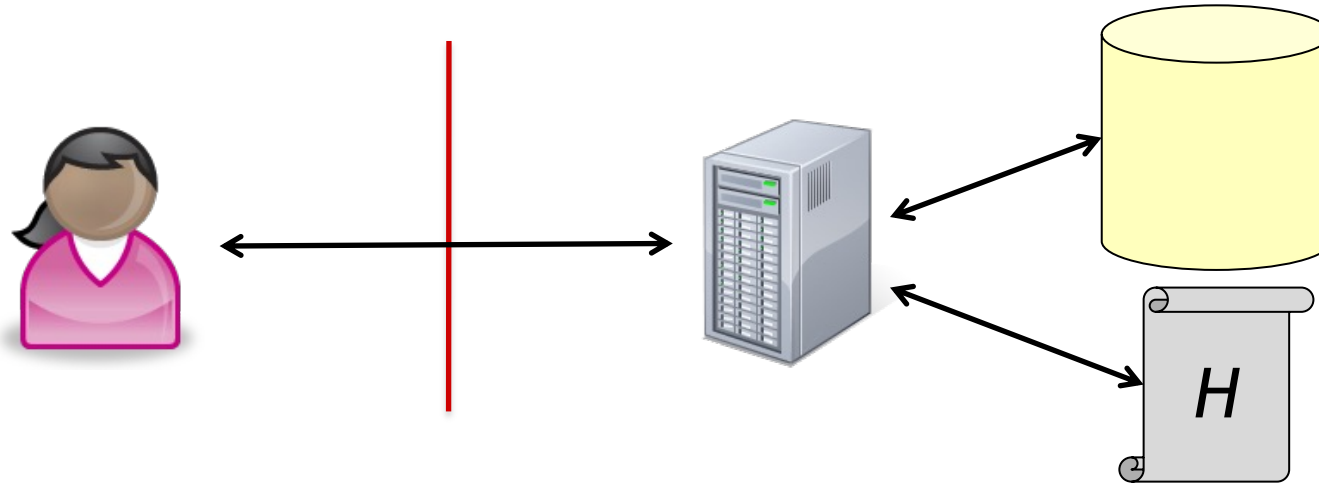
Data analysts often prefer this model of data release (**Why?**)

Common operations performed on data include:

- Stripping out names and other identifiers (**Suppression**)
- Grouping data values into less precise buckets (**Generalization**)
- Adding noise to records or groups of records (**Perturbation**)



# Mediated Query Processing



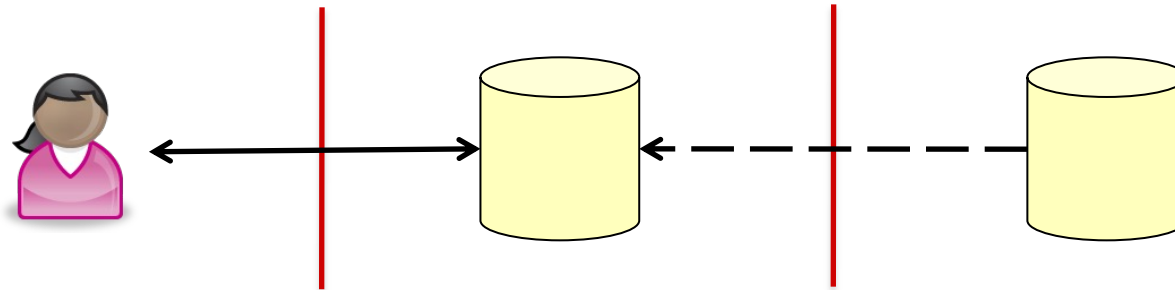
**Critical point:** Data is **not** released by the data owner

Since data is retained by the owner, they also retain control

- Do I think that this query is safe to answer?
- What other questions has this querier asked? Should this affect my answer?
- What type of perturbation is needed to make answering this question safe?

This data model is used by the US Census Bureau

# Outsourced Data Hosting



**Scenario:** Let's pay someone else to host our data

- Became popular with the increased prevalence of the web
- Increasingly interesting as cloud computing becomes a reality

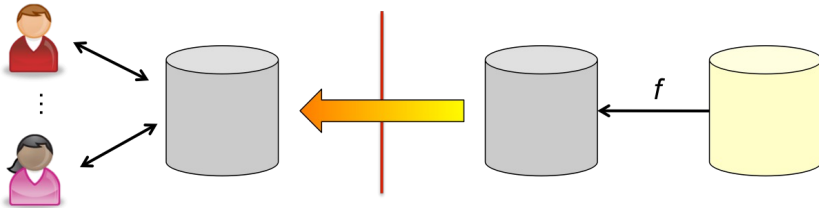
Potential uses include **offsite backups** and **outsourced DB management**

Depending on the reasons behind outsourcing, a variety of questions deserve some attention:

- Should the data host be able to read the data?
- Should the data host be able to learn about the organization of the data?
- Should queries be revealed to the data host?
- Is the data that is claimed to be hosted actually available?

This is an active area of academic research

# Each of these data models has various pros and cons

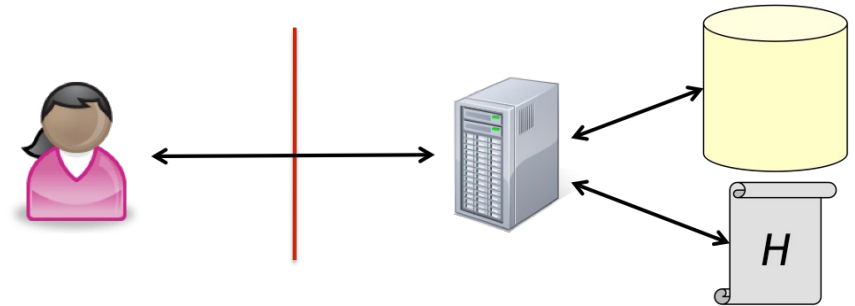


## Strengths

- Analysts get (nearly) complete access to data
  - Can explore data in novel/unpredicted ways
  - Can ask any questions they want
- Providers do **not** need to host data locally

## Weaknesses

- When is data “safe” to release?
- Balancing privacy versus utility?
- Quantifying anonymization?



## Strengths

- Analysts can ask many types of queries to the data store
- Providers can see all access to data and can adjust as needed

## Weaknesses

- Potentially need to store LOTS of query history
- How should data be perturbed?

# *Case Study:* k-Anonymity

L. Sweeney, “k-anonymity: A Model for Protecting Privacy,” International Journal on Uncertainty, Fuzziness and Knowledge-based Systems, 10 (5), 2002; 557-570.

# The state of the art for protecting privacy in the early 1990s was simply removing “identifiers”

<i>Name</i>	<i>Race</i>	<i>Birth</i>	<i>Gender</i>	<i>ZIP</i>	<i>Problem</i>
Aaron	Black	1965	M	02145	Short breath
Bob	Black	1965	M	02143	Chest pain
Christina	Black	1965	F	02133	Hypertension
Danielle	Black	1965	F	02137	Hypertension
Eve	Black	1964	F	02137	Obesity
Francine	Black	1964	F	02134	Chest pain
George	White	1964	M	02138	Chest pain
Harry	White	1964	M	02138	Obesity
Ian	White	1964	M	02134	Short breath
James	White	1967	M	02133	Chest pain
Kevin	White	1967	M	02133	Chest pain

**Question:** Who can see a problem with this?

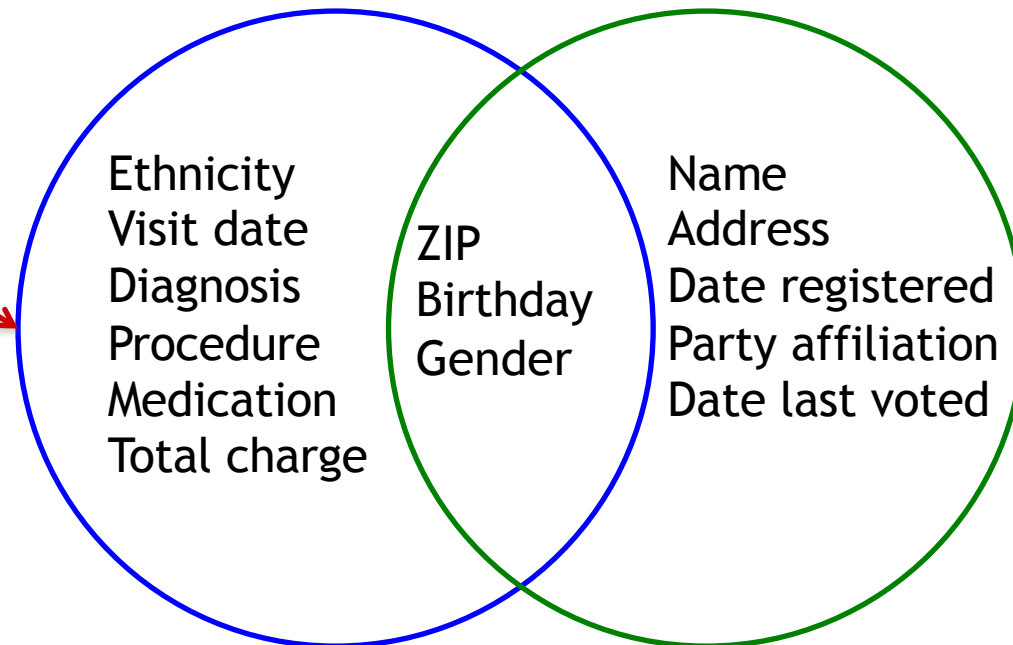
# Answer: Your name is not your only unique identifier!

**One example:** The triple (City, Birthday, Gender) is a unique identifier for 53% of the population, while (County, Birthday, Gender) identifies 18%

**Interesting attack:** Reidentifying medical records

Massachusetts Group Insurance  
Commission data set. Released to  
researchers and sold to industry.

Voter records. Cost: \$20.



*After joining these two datasets, Sweeny was able to recover the medical records of William Weld, the (then) governor of Massachusetts!*

# Some terminology...

The diagram illustrates the classification of data attributes in a table. Red arrows point from labels to specific columns: 'Explicit identifier' points to 'Name'; 'Quasi-identifiers' points to a bracket spanning 'Race', 'Birth', 'Gender', and 'ZIP'; 'Sensitive attribute(s)' points to 'Problem'. The table contains two rows of data: Aaron and Bob.

<i>Name</i>	<i>Race</i>	<i>Birth</i>	<i>Gender</i>	<i>ZIP</i>	<i>Problem</i>
Aaron	Black	1965	M	02145	Short breath
Bob	Black	1965	M	02143	Chest pain

...

Steps for “anonymize and release” data processing:

1. Remove all explicit identifiers
2. Manipulate rows to ensure that quasi-identifiers **cannot** be used to map specific individuals to sensitive attributes

Seems easy, right?

# k-Anonymity was one of the first rigorously studied anonymization methods

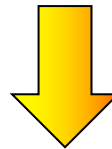
**High-level goal:** Each unique quasi-identifier should appear at least  $k$  times in the released data set

How can we accomplish this goal?

- Attribute generalization
- Attribute suppression
- Attribute perturbation

This provides a sort of plausible deniability...

<i>Race</i>	<i>Birth</i>	<i>Gender</i>	<i>ZIP</i>	<i>Problem</i>
Black	1965	M	02145	Short breath



Black	1967	*	021**	Short breath
-------	------	---	-------	--------------

*Let's see an example...*




# This is a 2-anonymous version of our hospital data table example

	<i>Race</i>	<i>Birth</i>	<i>Gender</i>	<i>ZIP</i>	<i>Problem</i>
→	Black	1965	M	0214*	Short breath
→	Black	1965	M	0214*	Chest pain
→	Black	1965	F	0213*	Hypertension
→	Black	1965	F	0213*	Hypertension
→	Black	1964	F	0213*	Obesity
→	Black	1964	F	0213*	Chest pain
→	White	1964	M	0213*	Chest pain
→	White	1964	M	0213*	Obesity
→	White	1964	M	0213*	Short breath
→	White	1967	M	0213*	Chest pain
→	White	1967	M	0213*	Chest pain

**Question:** Why is this table 2-anonymous?

- Each quasi-identifier appears (at least) 2 times

*Question:* Is the following table 3-anonymous?



<i>Race</i>	<i>Birth</i>	<i>Gender</i>	<i>ZIP</i>	<i>Problem</i>
Black	1965	M	0214*	Short breath
Black	1965	M	0214*	Chest pain
Black	1965	M	0214*	Hypertension
Black	1965	F	0213*	Hypertension
Black	1964	F	0213*	Obesity
Black	1964	F	0213*	Chest pain
White	1964	M	0213*	Chest pain
White	1964	M	0213*	Obesity
White	1964	M	0213*	Short breath

# k-Anonymity sounds great! So the data anonymization problem is solved, right?

## Problem 1: Solving the anonymization quality/efficiency trade-off

This table is 5-anonymous, but useless!



<i>Race</i>	<i>Birth</i>	<i>Gender</i>	<i>ZIP</i>	<i>Problem</i>
*	19**	*	*	Short breath
*	19**	*	*	Chest pain
*	19**	*	*	Hypertension
*	19**	*	*	Hypertension
*	19**	*	*	Obesity

**Question:** How can we define the “goodness” of a dataset?

- Less suppression/generalization/perturbation → better quality

**Fact:** Finding an optimal k-anonymization is an NP-Hard problem

Fortunately, heuristic methods do a pretty good job of this with fairly low overheads (see work by LeFevre et al.)

# Efficiency is solved, but what other problems are there?

**Problem 2:** How do we choose the value of  $k$  to use?

Essentially, there is no good answer to this question...

- How much better is 3-anonymity than 2-anonymity?
- Is the same value of  $k$  reasonable for all individuals in the dataset?
- How much does adjusting  $k$  impact the quality of the released data?



# More problems still...

**Scenario:** Bob has a record in the dataset, was born in the 1960s, and lives in the 15260 ZIP code

<i>Race</i>	<i>Birth</i>	<i>Gender</i>	<i>ZIP</i>	<i>Problem</i>
Black	196*	M	15260	Brain cancer
Black	196*	M	15260	Brain cancer
Black	196*	M	15260	Brain cancer
Black	196*	M	15260	Brain cancer

This is a 4-anonymous table, but Bob has brain cancer...

<i>Race</i>	<i>Birth</i>	<i>Gender</i>	<i>ZIP</i>	<i>Problem</i>
Black	196*	M	15260	Brain cancer
Black	196*	M	15260	Lung cancer
Black	196*	M	15260	Leukemia
Black	196*	M	15260	Bone cancer

This is a 4-anonymous table, but Bob has cancer...

# A generalization of this problem...

These problems emerge because the data set was not **diverse**

- All entries for a quasi-identifier map to the same sensitive attribute
- All entries for a quasi-identifier map to related sensitive attributes

Follow-on work addresses this but is subject to attacks of its own!

---

The bigger problem is that this class of solutions does not adequately model the knowledge of the attacker

- I know that Bob visited the hospital and should be in this data set
- I know that Bob has some type of cancer
- ...

Recent work on **differential privacy** works for **any** attacker, but uses the mediated query model

In short, this is still a very active research area

# Data Privacy Summary

So we talked about three types of models for managing private data

- Anonymize and release
- Mediated query processing
- Outsourced data hosting

k-Anonymity is one solution in the “anonymize and release” model

- Strip out explicit identifiers
- Be sure that each quasi-identifier appears at least  $k$  times

How do we manage diversity and model attacker knowledge?

- Recent work does this with limited success

**Take away point:** Data anonymization is hard. “Anonymization” is probably a flawed term, as it is hard to quantify...

**Next:** Blockchains

# Blockchains

What *is* a blockchain?

Building a blockchain

- Important properties
- Crypto crash course
- Basic construction

Cryptocurrency: Bitcoin as an example

Other application areas



# Blockchains

What *is* a blockchain?

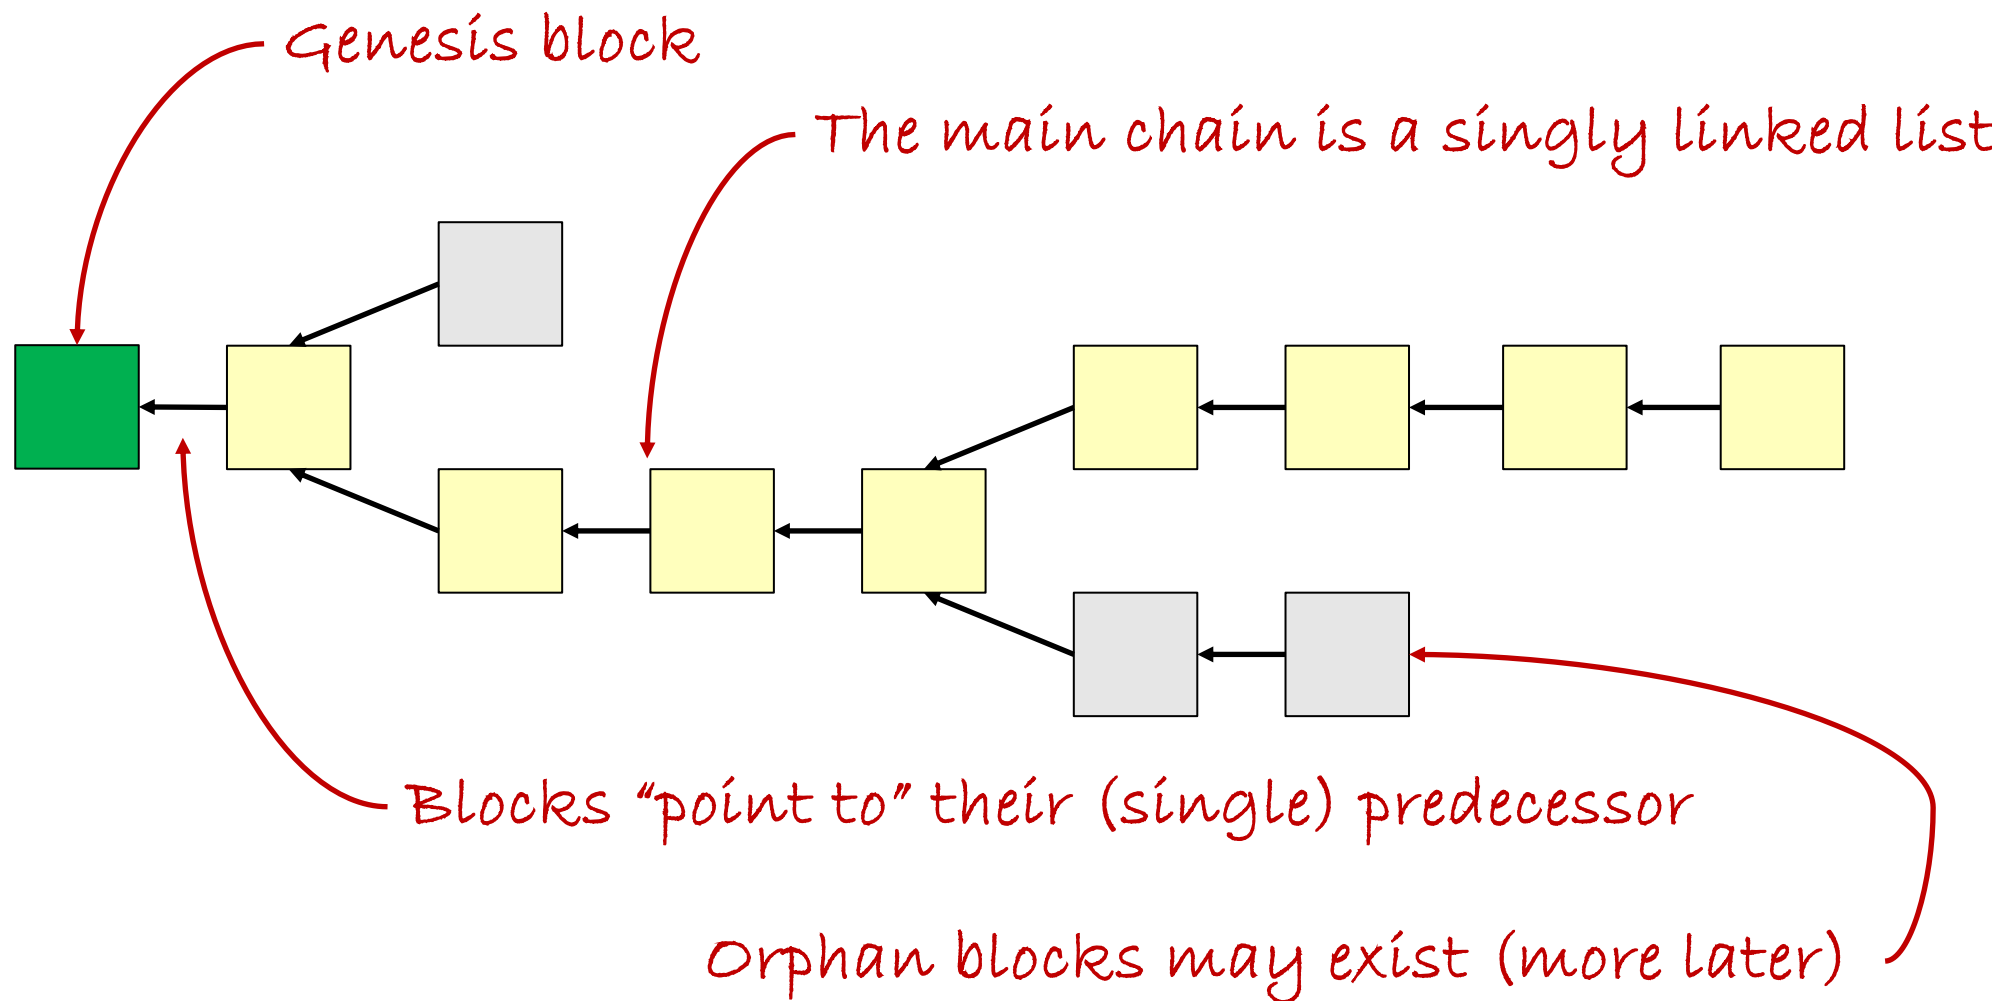
Building a blockchain

- Important properties
- Crypto crash course
- Basic construction

Cryptocurrency: Bitcoin as an example

Other application areas

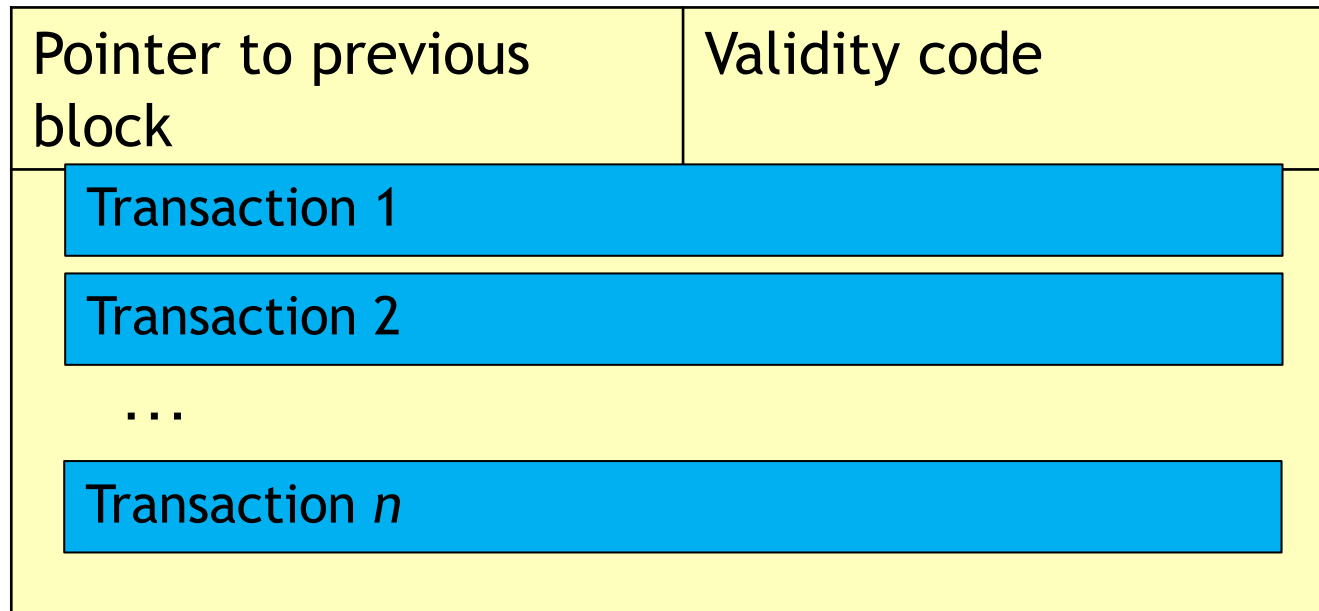
# A blockchain is just a distributed data structure



---

*What's in a block? How is the blockchain extended?*

# Each block can be thought of as a page in a (decentralized) transaction registry



What's a transaction? Depends upon the application...

- A transfer of currency (e.g., Bitcoin)
- Events in distributed monitoring application
- Claims filed against an insurance policy
- ...

*Yes, I'm waving my hands here...*

The validity code is used for obtaining consensus and making the block (and blockchain) tamper evident

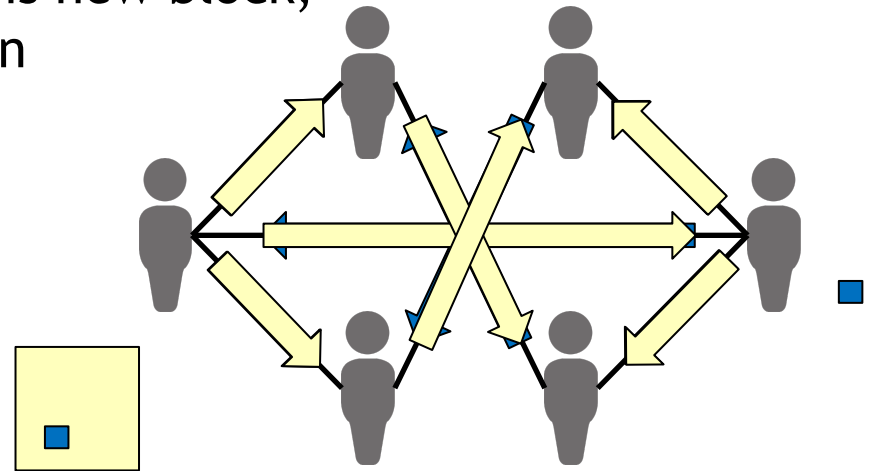
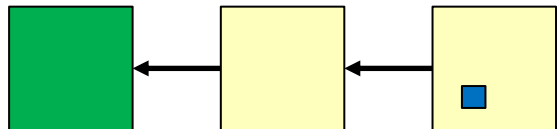
# How is the blockchain extended?

Blockchain protocols are based on decentralized (P2P) networks

Extending the blockchain

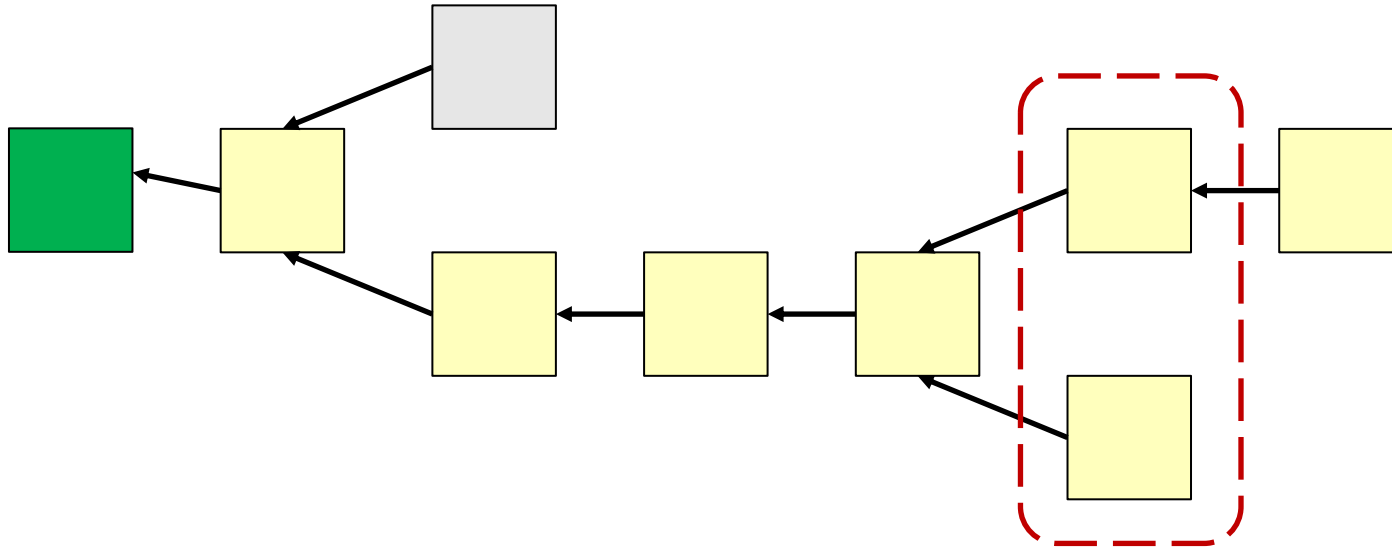
1. Individual transactions are **flooded** through the network
2. If valid\*, transactions are added to in-progress blocks
3. Participants **solve a puzzle** to compute the block validity code
4. New blocks are flooded through the network
5. If valid\*, the peers build upon this new block, thereby extending the blockchain

*"Mining"*



# Remember orphan blocks?

Nodes build upon the longest chain that they know



What happens if blocks are created concurrently?

- Build upon the first received

The result? Blocks can become orphaned.

- What happens to transactions in orphaned blocks? Flood again!
- When can a transaction be considered committed?

# Further considerations

If properly constructed, a blockchain can provide a **highly-available, publicly-verifiable, immutable, consensus-based** distributed transaction register

Technical and application-specific questions:

- How do we manage identities in a decentralized system?
- How can we efficiently verify transaction validity? Block validity?
- How do we ensure that the blockchain remains immutable?
- How can we balance throughput and consistency?
- How do we incentivize participation and mining?

Let's dig deeper and find out...

# Blockchains

What *is* a blockchain?

Building a blockchain

- Important properties
- Crypto crash course
- Basic construction

Cryptocurrency: Bitcoin as an example

Other application areas

# What properties should a minimally-interesting blockchain offer us?

A minimal set of important properties includes:

- **Availability** of the blockchain
- **Append-only** and **immutable** transaction log
- **Tamper-proof** and **authenticated** transaction records
- **Verifiability** of transactions and structural invariants
- **Consensus** on the transactions to be included in the blockchain

Let's build a blockchain structure that provides these properties

**Note:** Assume an **open system** in which anyone can participate



# Proof-of-work systems are computational puzzles that can be used to rate limit behavior

A **computational puzzle** is a problem that is *computationally-difficult* to solve, but *easy* to verify

We can build a simple computational puzzle with tunable hardness using hash functions:

- **Given:** a hash function  $H(\cdot)$ , a message  $m$ , and a hardness  $h$
- **Compute:** a value  $v$ , such that  $H(m \parallel v)$  ends with  $h$  0-bits

Why is this hard to solve?

- It is computationally-difficult to invert a hash function
- Must brute-force search for a  $v$  that satisfies the puzzle constraints
- I.e.,  $O(2^h)$  tries to solve the puzzle

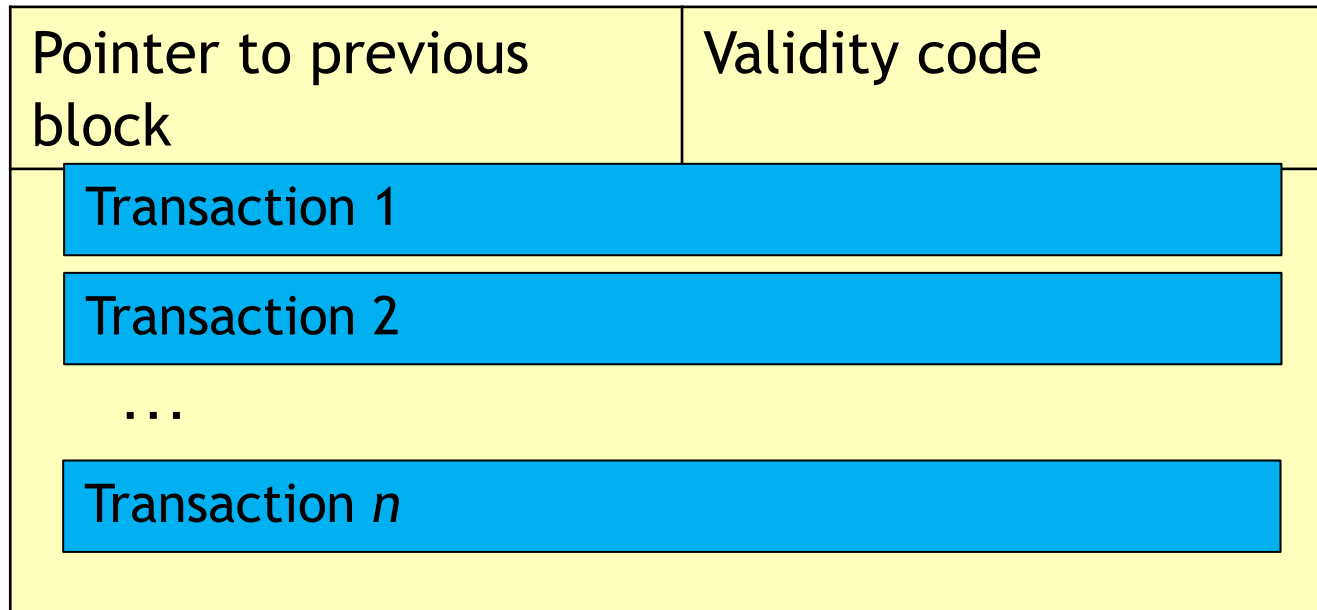
Preimage resistance



*Let's detail the structure of a basic blockchain*

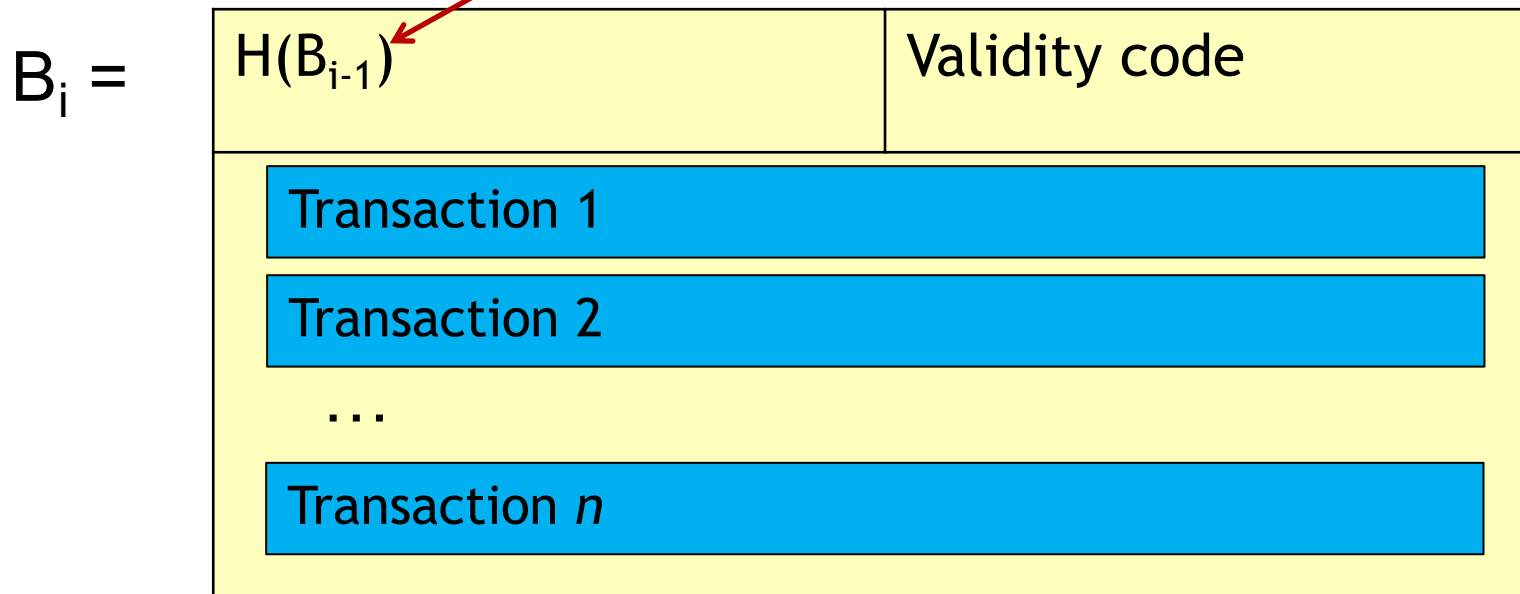
We won't change our basic block structure, we'll just fill in the details...

$B_i =$



# We won't change our basic block structure, we'll just fill in the details...

A block's address is just the hash of its contents



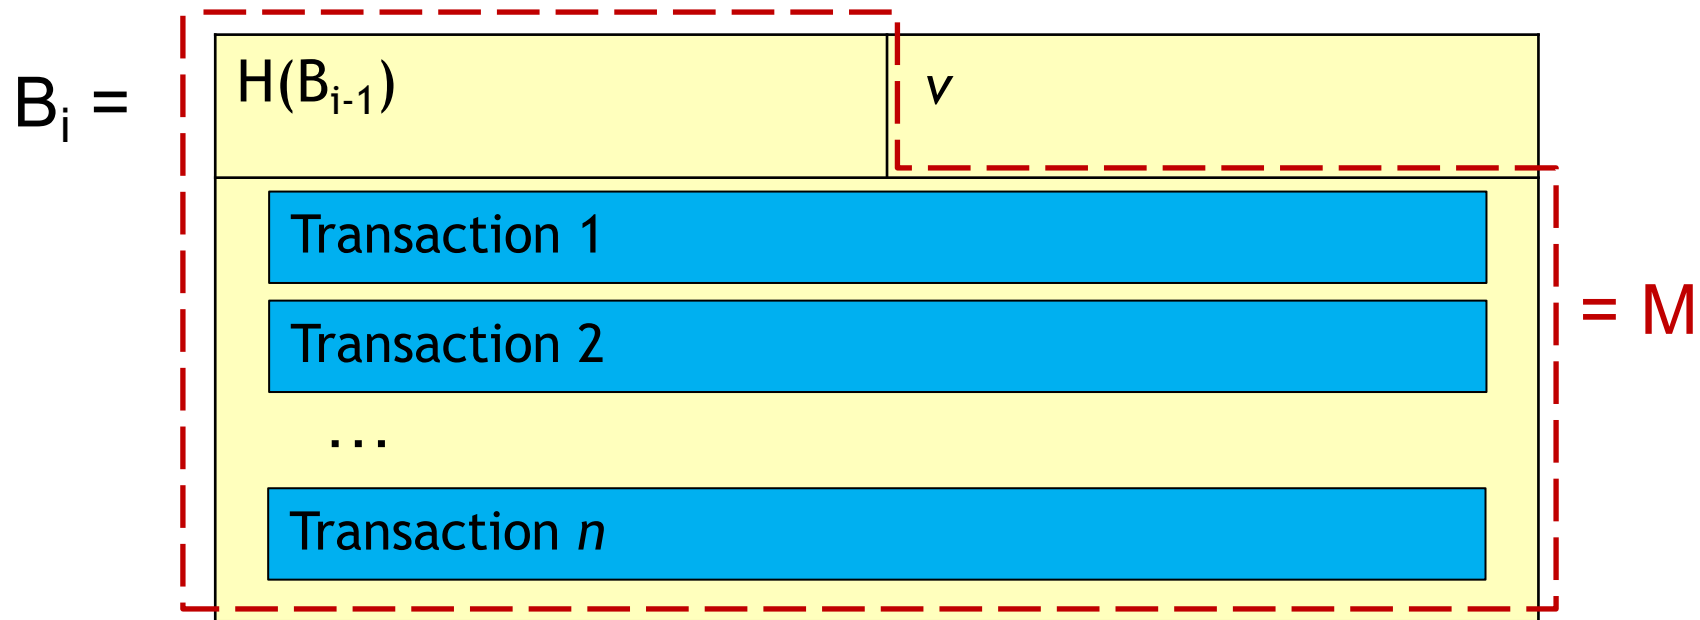
$H(B_{i-1})$  is a unique value identifying  $B_{i-1}$

- Why? **Collision resistance** and **2<sup>nd</sup> preimage resistance** of  $H(.)$

**Side-effect:** Changing block *contents* changes the block *address*

- This will be useful to us later...

We won't change our basic block structure, we'll just fill in the details...

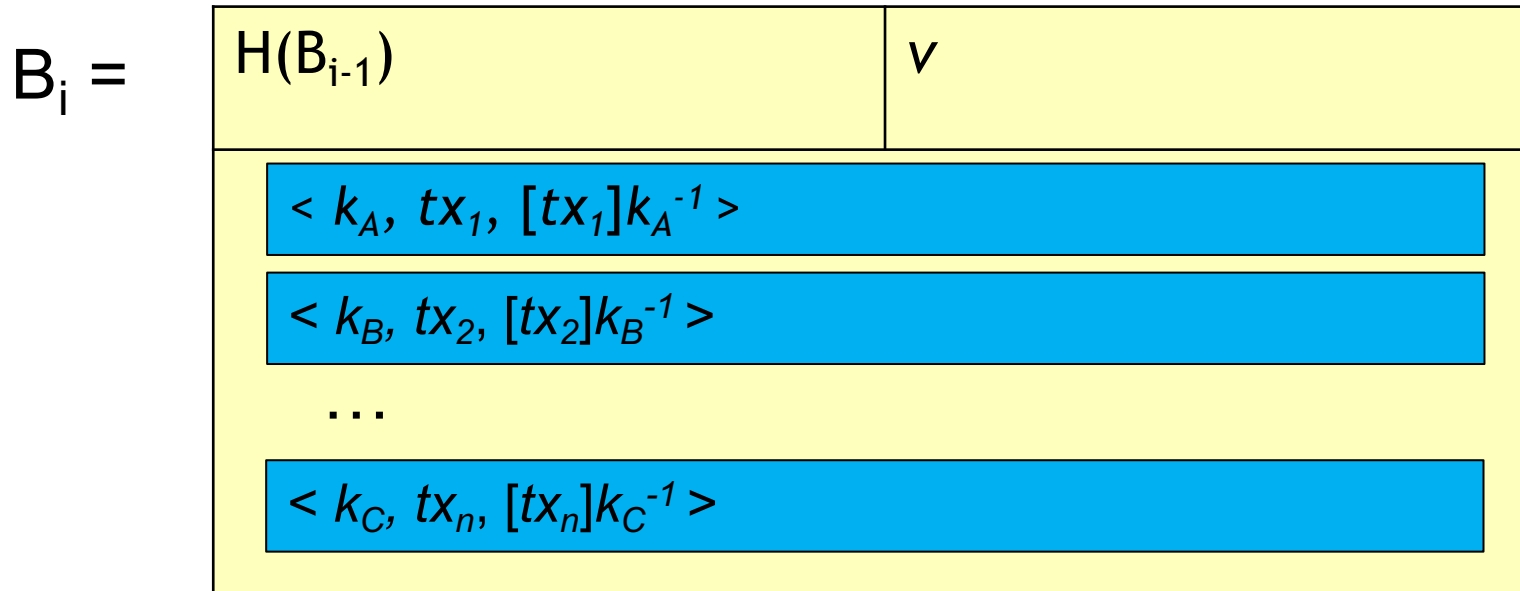


We'll use a hash puzzle to compute the validity code

- i.e., find  $v$  s.t.  $H(\textcolor{red}{M} || v)$  ends with  $h$  0-bits

How do we represent individual transactions?

# Main idea: *Public* keys represent user identities



Each transaction record consists of three parts:

- A public key (e.g.,  $k_A$ ) ← Who registered this transaction?
- The transaction contents (e.g.,  $tx_1$ ) ← Application specific...
- A digital signature over  $tx$  (i.e.,  $[tx_1]k_A^{-1}$ ) ← Ensures accountability

---

*That's it! What about our properties?*

# What properties should a minimally-interesting blockchain offer us?

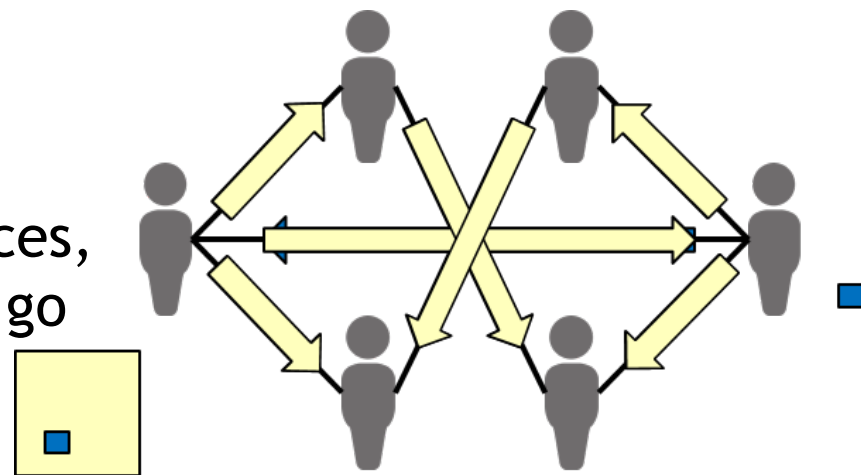
A minimal set of important properties includes:

- ✓ ● **Availability** of the blockchain
  - **Verifiability** of transactions and structural invariants
  - **Tamper-proof** and **authenticated** transaction records
  - **Append-only** and **immutable** transaction log
  - **Consensus** on the transactions to be included in the blockchain
- 

Gossip -> Availability

- Each transaction floods the network
- Each block floods the network

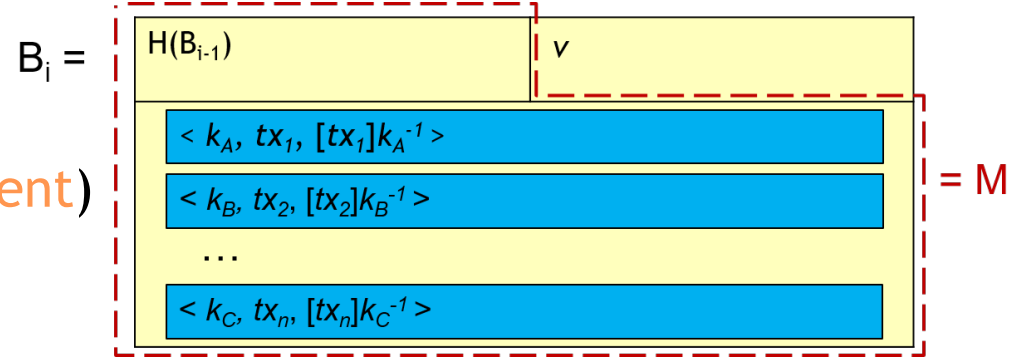
Copies of txns/blocks exist in **many** places, remaining available as nodes come and go



# Verifiability is more nuanced

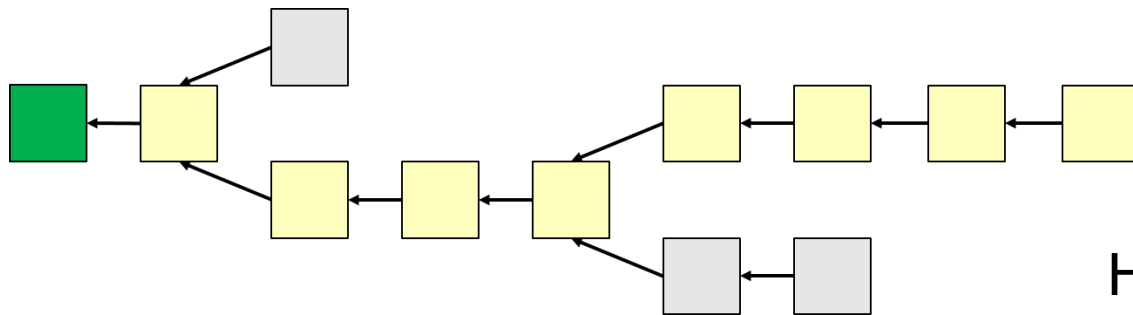
Two steps to validate each transaction

- Validate the txn signature (**easy**)
- Validate the txn (**application dependent**)



To validate the block

- Validate each transaction (**easy**)
- Check the validity code (**easy**: compute  $H(M || v)$  and check structure)



Validating structural invariants

- Is this new block valid? (**easy**)
- Does it build upon the most recent block I know of? (**easy**)

Handling forks isn't hard, either



# Digital signatures ensure that transactions are authenticated and tamper-proof

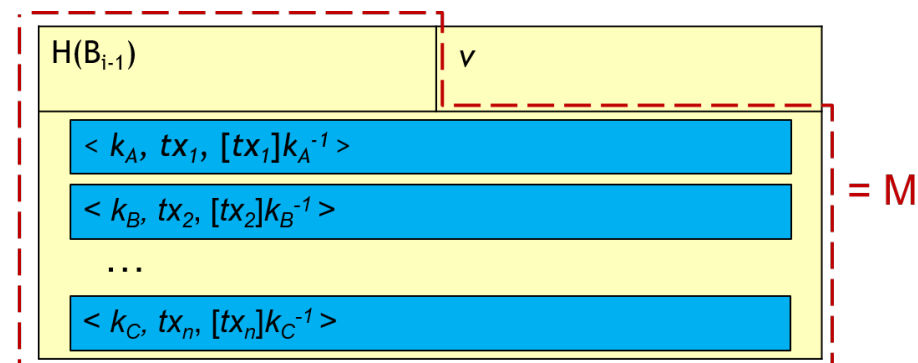
A minimal set of important properties includes:

- ✓ ● **Availability** of the blockchain
  - ✓ ● **Verifiability** of transactions and structural invariants
  - ✓ ● **Tamper-proof** and **authenticated** transaction records
  - **Append-only** and **immutable** transaction log
  - **Consensus** on the transactions to be included in the blockchain
- 

Authentication via non-repudiability

- $k_A$  identifies user A
- Only  $k_A^{-1}$  can produce  $[tx_1]k_A^{-1}$
- **Result:**  $tx_1$  was registered by A

$B_i =$



Transactions are also tamper-proof

- Changing a transaction invalidates its signature
- “Updating” a record will invalidate the proof-of-work  $v$ 
  - Why? 2<sup>nd</sup> preimage resistance of  $H(.)$

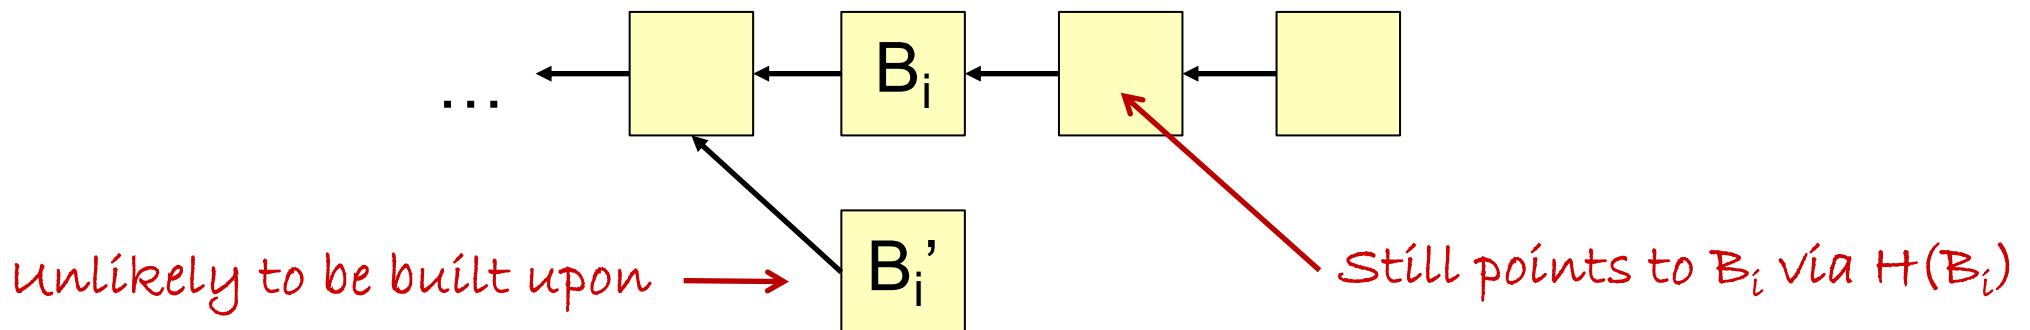
# Modifying a block creates a fork in the blockchain

A minimal set of important properties includes:

- ✓ ● **Availability** of the blockchain
  - ✓ ● **Verifiability** of transactions and structural invariants
  - ✓ ● **Tamper-proof** and **authenticated** transaction records
  - ✓ ● **Append-only** and **immutable** transaction log
  - **Consensus** on the transactions to be included in the blockchain
- 

**Recall:** *Deleting* a block is a non-starter

*Modifying* a block creates a fork in the chain since  $H(B_i) \neq H(B_i')$



# The blockchain is advanced through distributed work and community validation

A minimal set of important properties includes:

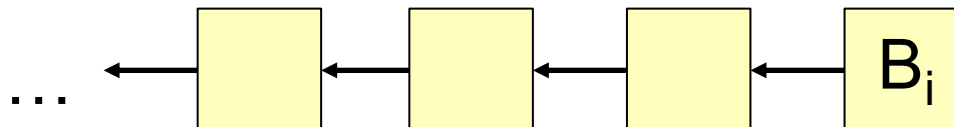
- ✓ ● **Availability** of the blockchain
  - ✓ ● **Verifiability** of transactions and structural invariants
  - ✓ ● **Tamper-proof** and **authenticated** transaction records
  - ✓ ● **Append-only** and **immutable** transaction log
  - ✓ ● **Consensus** on the transactions to be included in the blockchain
- 

No single node is in charge of advancing the blockchain

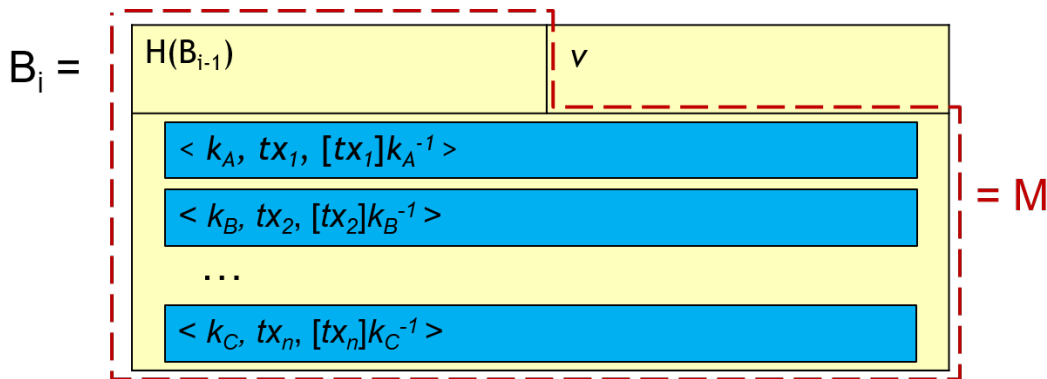
- Must solve proof of work to finalize new block,  $B_i$
- This is a random process governed by compute power

A majority of nodes must accept each new block

- Think of building off of a new block as “voting” for it



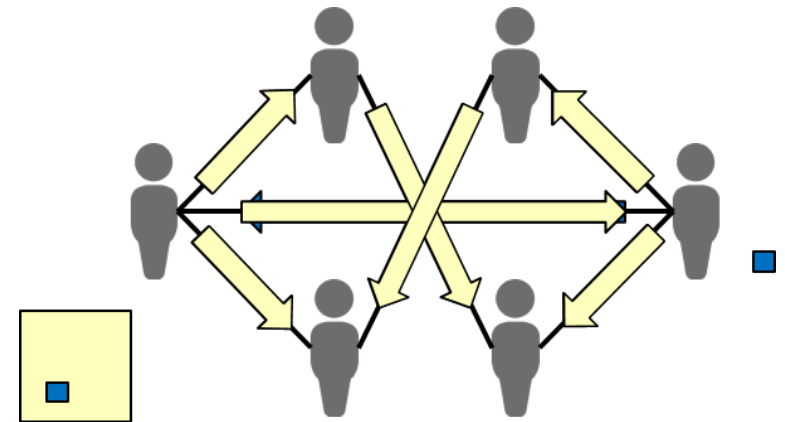
# This technical foundation is helpful, but there are still application-specific considerations



*How do we incentivize participation (e.g., mining?)*

*How should transactions be structured to make them efficiently verifiable?*

*How can we balance the efficiency of arriving at consensus with the scalability and/or openness of the network?*



# Blockchains

What *is* a blockchain?

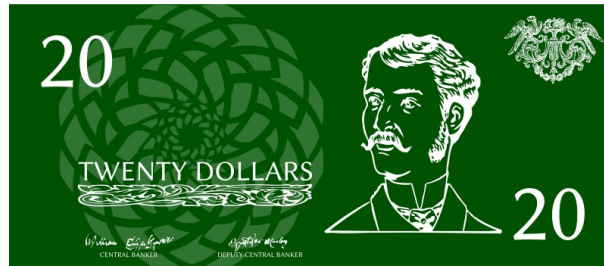
Building a blockchain

- Important properties
- Crypto crash course
- Basic construction

**Cryptocurrency: Bitcoin as an example**

Other application areas

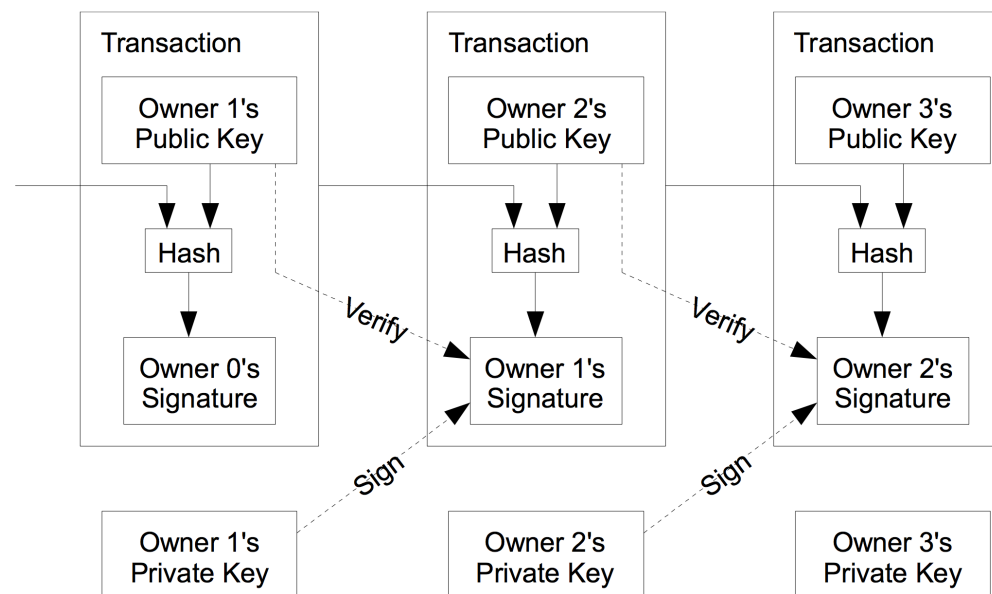
# The goals of any currency



- **Counterfeit detection:** How do I know this bill is authentic?
- **Double-spending protection:** How do I know that nobody can claim this money belongs to them and not to me?
- **Question:** Which is easier to accomplish with physical cash?
  - Counterfeits can be very sophisticated; evasion is a cat-and-mouse game
  - Anti-counterfeiting technologies aim to be expensive to duplicate
  - Double-spending is easy: whoever holds the bill is the owner!

# What could we do in a distributed, digital setting to achieve these goals?

- In a distributed system, agreeing on whether a coin is valid is a form of the Byzantine Generals' problem
  - Need to reach agreement without central authority
- Spending can be tracked by using public-key crypto and keeping a public, distributed, signed log of transactions relating to a coin



- How do we prevent **double-spending**?

# How can we reach a consensus about which transaction occurred first?

- Like in physical counterfeiting, we want it to be **overwhelmingly expensive** to cheat
- What about **voting**?
  - If a second transaction shows up before agreement, cheating is detected
  - How do we prevent duplicate votes?
- Bitcoin uses a **proof-of-work** system to rate-limit voting
- The version of history written by the group with the **most computing power** becomes the truth
  - As long as *most*\* users are honest, the log is correct!



# Bitcoin is a cryptocurrency based upon a blockchain like the one that we've been exploring today

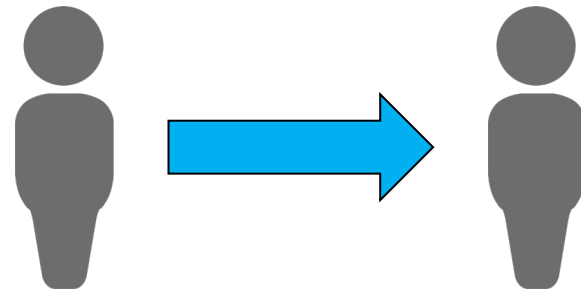
## Technical parameters:

- Digital signatures using ECDSA on the secp256k1 curve
  - ~128 bits of security
- Hashes computed using SHA-256
- Proof-of-work slightly different than the one we used
  - Can adjust with higher precision than 1 bit
  - Goal: ~10 minutes to produce a block

## Two classes of transactions:



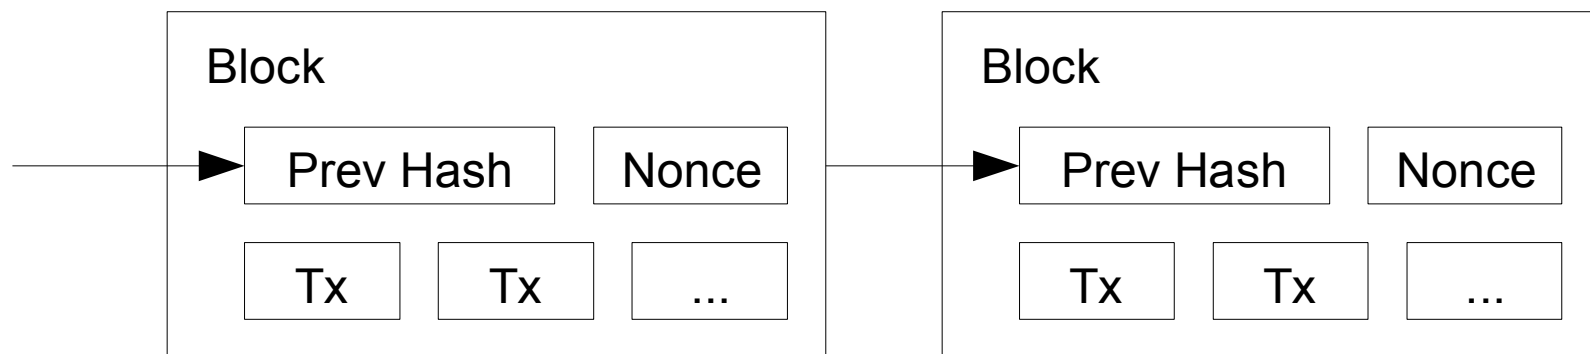
*Coinbase transactions*



*Transfers\**

# Bitcoin's Proof-of-work System

- Hashcash: Originally developed for email (**why?**)
- To send message  $M$ , Alice must compute **nonce**  $n$  such that  $H(M||n)$  starts with  $b$  '0' bits
  - How long will this take if  $H$  is a **cryptographic hash function**?
- Bitcoin uses a **block chain** of transaction blocks, where each one requires a Hashcash nonce: longest chain wins!



**One vote per CPU/FLOP**

# So how do I know I've received money for a product or service?

- Just because a transaction is in *some* block, doesn't mean that this block will be chosen by the majority
  - Blocks that aren't incorporated into the "canonical" chain are called **orphaned blocks**
- In order to accept Bitcoin, a vendor must wait until a transaction has sufficient blocks built on top of it
  - Otherwise, the spender may be able to **double-spend**
- Unfortunately, this can take a long time
  - Blocks created every ~10 minutes (tuned with  $b$  automatically)
  - Block restricted to **1 MB**
- **Bitcoin XT** is a fork of the Bitcoin software
  - Allows clients to vote on whether to increase the block size
  - Once 75% of clients vote yes, and after a 4-week waiting period, 2 MB blocks will be accepted
  - This would create a fork in the block chain

# The steps so far...

1. **Broadcast** new transactions to all nodes
2. Each node **collects** new transactions into a block
3. Each node searches for **nonce** to complete block
4. First successful node **broadcasts** complete block
5. All nodes **verify** all transactions in the new block
6. Nodes **attest** to a block's correctness by building upon it for the next block

# How does bitcoin incentivize mining?

## Incentive 1: Mining rewards

- The first transaction in a block is a coinbase transaction
- Initially, 50 BTC per block, halves every 210,000 blocks
  - Current reward: 6.25 BTC
  - Maximum currency volume: 21,000,000 BTC (currently ~19.5M)

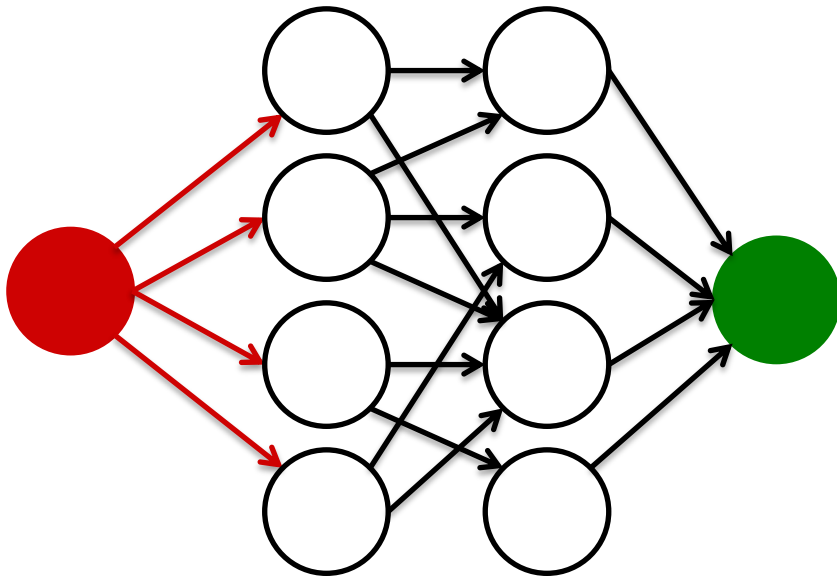
## Incentive 2: Transaction fees

- Small fees attached to each transaction
- Fees claimable by the node that creates a block

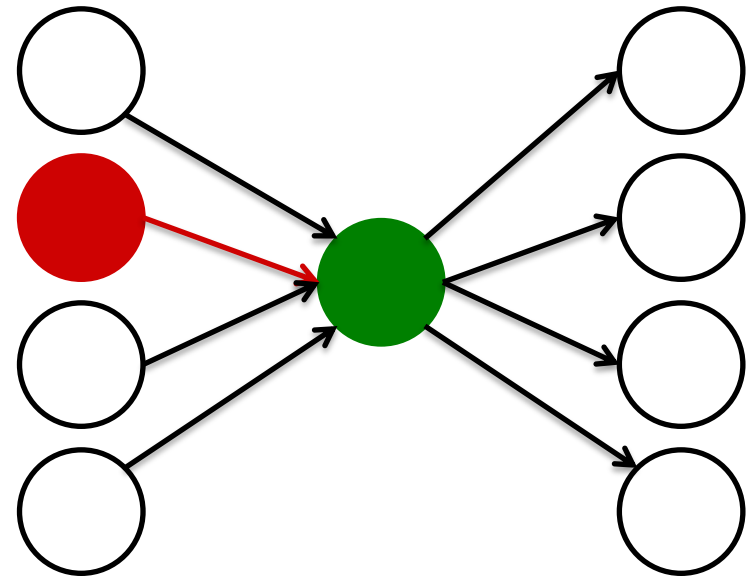
**Mining pools:** groups of nodes to work together, split rewards

# Is Bitcoin really as private as cash?

- Since all transactions are public, looking at the history can reveal trends
- **Idea:** Transfer coins between multiple addresses



Fork-merge



Mixing

# BlockChain Concluding thoughts

At its core, a blockchain is simply a data structure

Blockchains have a number of useful features

- **Distributed** and **highly-available** transaction ledger
- **Append-only** and **immutable** record of activities
- **Tamper-proof** and **authenticated** transactions
- (Perhaps public) **verifiability** of transactions and structure
- **Consensus** required to extend the structure

Applications to many interesting classes of problems

Flexibility can be extended via the use of **smart contracts**

# Cryptocurrency Conclusions

Bitcoin and other cryptocurrencies satisfy properties to be a usable **currency**

- Counterfeit detection
- Double-spending resistance

Although often thought of as “the digital equivalent of cash,” Bitcoin is **not** entirely private

- Public ledger, statistical analysis

**Proof-of-work systems** enable solution to Byzantine fault tolerance

Proposals such as **Zerocash** expand Bitcoin to strengthen its guarantees



# Topic Area 1: The Basics

## *Objectives:*

- Why study computer security? Is it really important?
- What are the CIA properties?
- How do these properties parameterize most computer security goals?
- Security is a relative (**not** absolute!) property

# Topic Area 2: The Tools

## *Objectives:*

- Understanding cryptography as a black box
  - When should I use public key versus secret key cryptography?
  - When are block ciphers good to use? Stream ciphers?
  - Which modes of operation are good for which types of task?
  - Why are hash functions the “duct tape” of security protocols?
- Understanding cryptography in depth
  - What is semantic security? Why is it useful?
  - How do ciphers like AES work? Why are they safe? Why are they fast?
  - Why is RSA secure?
  - How does the mathematics behind RSA influence attacks against implementations?

# Topic Area 3: Assembling the Primitives

## *Objectives:*

- I have a toolbox full of cryptography, how do I use it?
- Basic cryptographic protocols
  - Identifying/authenticating users
  - Mutually authentication with services
  - Safely exchanging secrets
  - Managing public keys in a realistic manner
- Subtleties galore...
  - Why doesn't "just encrypt it" or "just HMAC it" work?
  - Protecting against replay/reflection attacks
  - Whose clock?
- **Main point:** Having tools is important. Using tools properly is more important.

# Topic Area 4: Security Challenges in the Real World

## *Objectives:*

- Understanding security in the abstract is good, but how do we put it into practice?
- What issues are involved with OS/application protection?
  - Integrity of mechanism
  - Storing, managing, accessing, and updating user permissions
  - Insiders vs. outsiders
- Many real-world attacks come from violated design assumptions
- How is privacy different from confidentiality? How is it the same?

# So, what was the point of the term project?

## Phase 1:

- Security is not an absolute property!
- How will we define “secure” for this application domain?
- What design principles do we need to respect?
- **Writing:** Informal writing can facilitate brainstorming

## Phase 2:

- Build a non-trivial code base in a cooperative manner
  - Team work is the norm outside of classes
  - Learn about version control and useful tools
- Think about functionality, not security
- Get used to documenting the code that you produce

# So, what was the point of the term project?

## Phases 3 and 4:

- Adding security to a non-trivial codebase as an afterthought is a pain...
  - So think about security from the get-go!
- Given the tools that we've learned about in class, how can we apply them?
- Given good ideas, implementing them still involves lots of subtlety
- **Writing:**
  - Informal specifications can guide development and facilitate teamwork
  - Precision, precision, precision!

## Phase 5:

- Threat models are just assumptions: what happens when they change?
- When securing a system, it is helpful to think like an attacker
- **Writing:** Formal writing
  - Good vulnerability documentation is a must
  - Good system designs help with formal verification **and** software engineering

# What do I hope that you have learned?

## *Academic/Professional*



- What does it mean for a system to be secure?
- Develop a **working knowledge** of cryptography
  - How to use properly, rather than how to design/prove/implement
  - I.e., build systems using tools, not craft new tools per se
- Understand the subtleties of protocol/system design
- Get comfortable assessing the security of non-trivial software systems

## *Everyday/Real World*

- What types of security/privacy issues should you be aware of?
  - Private email
  - Exposed network traffic
  - Your data in public
- Realistically evaluating the security of services that you use
- Managing the complexities of non-trivial team work

